

Onwards: The Formal Tradition Was Waiting for Its Executor

Onwards: The Formal Tradition Was Waiting for Its Executor

Juan Carlos Ghiringhelli Pragmaworks · May 2026

Abstract

The formal tradition of correct computing — from Aristotle’s categories through Hoare’s contracts to Honda’s session types — produced 2,376 years of theorems that proved how programs could be made correct. Very little of it reached widespread production use, and the parts that did (type systems in modern compilers, contract assertions in some safety-critical pipelines) reached far less of it than the theorems would have predicted. This essay argues the gap was structural, not intellectual: the executor required to sustain formal discipline across teams, deadlines, and rotations was the human engineer, and the cost of carrying every theorem in memory while applying each correctly to evolving code was never recoverable. In 2017, the transformer architecture produced a different kind of reader — one that draws on a substantial slice of the formal corpus without the fatigue that erodes human application of it, and benefits immediately from the correctness it generates. I argue that Generative Specification (GS) is the discipline that connects them: it activates formal theory through specification rather than annotation, addresses the three structural causes of abandonment (fatigue, single-target economics, fragmentation), and exhibits the functional structure of biological autopoiesis as a consequence of being a self-maintaining formal system. The essay traces the full lineage, identifies the human author as the *Nous* (the intent, the *what* and *why*), the body of formal theory as the *Logos* (the laws governing correct work, the *how*), and the AI executor as the worker who operates under the *Logos* without the erosion human practitioners suffer. It names the structural category that distinguishes this from AGI: *directed formal autopoiesis* — a self-creating system that converges toward mathematical correctness through formally grounded increments. Loom, an AI-native language compiling to five targets from a single source, is offered as the language-layer demonstration backing the claim.

Keywords: programming paradigms, generative specification, formal methods, type systems, AI-assisted development, biological isomorphisms, autopoiesis, language-oriented programming

“What we cannot speak about clearly, we must pass over in silence.” — Wittgenstein, Tractatus Logico-Philosophicus (1921).

Much of what we used to pass over can now be said directly.

I. The Question

For as long as people have committed rules to writing — from the Code of Hammurabi nearly four millennia ago through the surviving philosophical treatises of antiquity to the present — they have asked the same question: *can meaning be expressed precisely enough that correct behavior can be mechanically derived from it?*

The Babylonians encoded legal contracts in cuneiform so judges could derive verdicts mechanically. The Greeks developed syllogistic logic so philosophers could derive truth from premises. The medieval scholastics built inference engines to derive theology from axioms. In 1666, Leibniz proposed the *Characteristica Universalis* [4] — a universal formal language in which all human knowledge could be expressed — combined with a *calculus ratiocinator*, a mechanical reasoner that would derive correct answers from the specification alone. “*If controversies were to arise,*” he wrote, “*there would be no more need of disputation between two philosophers than between two accountants. For it would suffice to take their pencils in their hands, sit down to their slates, and say to each other: Let us calculate.*”

In 1936, Church [9] and Turing [10] proved that computation itself is formalizable. In 1969, Hoare proved that programs can carry mathematical contracts [12]. In 1978, Milner proved that types are propositions and programs are proofs [16]. In 1988, Meyer made contracts executable syntax [19]. In 2000, Fielding proved that a stateless system can navigate itself through self-describing responses alone [24].

Every one of these was correct. Every one was published. Few reached the scale of adoption the theorems would have predicted.

This essay is about why, and about what changed.

II. The Formal Tradition

The intellectual lineage of correct computing, in the form we can read from the surviving record, runs 2,376 years. Older traditions almost certainly existed — every generation uncovers fragments of advanced mathematical and logical work from cultures earlier than the canonical Greek sources — but the continuous, written, peer-checkable arc begins with Aristotle’s *Categories* (350 BCE) [1], the first surviving attempt to classify things into kinds such that only

certain operations are valid on certain kinds. The first type system, not by metaphor but by definition. It passes through Euclid's axiomatic method (300 BCE) [2]: start with axioms, apply rules, arrive at truth — which is require:/ensure: with a different notation and the same semantics. It reaches Leibniz's dream (1666) [4], the specification-as-derivation-engine. It enters the modern era with Boole's algebra of logic (1854) [5], Frege's predicate calculus (1879) [6], Russell's type theory to repair Frege's paradox (1910) [7], Gödel's incompleteness to set the ceiling (1931) [8], Church's lambda calculus and Turing's machines to formalize computation itself (1936) [9, 10], and Curry and Howard's correspondence to prove that propositions *are* types and proofs *are* programs (1934–1969) [11].

Then the decisive decade. Hoare gives us program correctness as mathematical contracts (1969) [12]. Dijkstra gives us weakest preconditions and the most important sentence in the history of software engineering: “*Program testing can be used to show the presence of bugs, but never their absence*” (1976) [14]. Denning builds the information flow lattice — the proof that security labels form an algebraic structure and information must flow only from lower to higher clearance without explicit declassification (1976) [15]. Milner derives Hindley-Milner type inference, allowing the compiler to fill in what annotations omit (1978) [16].

The arc continues. Girard proves linear logic: some resources must be consumed exactly once (1987) [18]. Meyer makes design by contract a programming language feature (1988) [19]. Honda invents session types: communication protocols verified at compile time (1993) [21]. Kennedy formalizes units of measure (1996, in his Cambridge thesis on a Standard ML-based system) [22] — the dimensional-type discipline later realized in F#'s unit system, so that `Float<usd>` and `Float<eur>` are distinct types and cannot be accidentally added. Myers and Liskov build JIF, the first working implementation of Denning's 1976 lattice in a production compiler (1997) [23]. Plotkin and Power formalize algebraic effects (2001) [25]. O'Hearn develops separation logic for local reasoning about memory (2002) [26]. Kephart and Chess define MAPE-K, the feedback loop for self-adaptive systems (2003) [27]. Dwork defines differential privacy mathematically: $\text{@dp}(\epsilon=0.1)$ as a type-level annotation (2006) [30]. Honda and Yoshida extend session types to multiparty choreography (2008) [31]. Shapiro proves CRDTs — data structures whose merge operations are algebraically guaranteed to converge (2011) [33].

Every idea was published. Every proof was sound. The arc runs 2,376 years. The question is not whether the tradition was correct. The question is why so little of it reached production at the scale its theorems would have predicted.

The answer is three causes, and they are precise.

Annotation fatigue. The formal annotations these systems require are correct but expensive for humans to write, impossible to maintain as code evolves, and culturally foreign to most working engineers. Myers and Liskov's JIF shipped in 1997. It proved that Denning's 1976 lattice theory works in a real compiler. Nobody used it. Annotating a million-line codebase with security labels, by hand, indefinitely, across changing teams and deadlines, is not a trade-off working engineers make. The theory was never wrong. The cost of applying it was never recoverable.

Single-target value economics. Adding a unit type to Python is not worth the cost for Python alone. The annotation pays for itself only when it generates output across multiple targets simultaneously — a Rust newtype, a TypeScript branded type, a JSON Schema extension, an OpenAPI field. One annotation, five targets, five times the value. Without multi-target compilation, the annotation cost was always greater than the single-target benefit.

Tooling fragmentation. Type theory lives in compilers. Security labels live in audits. SLOs live in dashboards. Deployment configs live in YAML. Privacy obligations live in legal documents. They never meet. They never meet because connecting them would require maintaining five separate systems, and nobody builds integrations between a type checker and a Kubernetes manifest because the abstractions live at different layers of a stack that was never designed to be unified. The theories were correct. The tools were scattered. Nobody could carry the full formal tradition in one place because no place existed to carry it.

These three causes are not cultural failures or failures of will. They are structural constraints on what human practitioners can sustain. The formal tradition was not abandoned because it was wrong. It was abandoned because the available executor — the human engineer, working under deadlines, carrying the theories in memory, applying them by hand, maintaining the annotations indefinitely — could not hold it all without degrading it.

The three causes above — fatigue, single-target economics, fragmentation — are structural facts about the artefact (the theories and their tooling). At the layer below them sit three parallel constraints about the *human practitioner* who was supposed to carry the artefact through working life:

- **Learning.** No practitioner career is long enough to achieve simultaneous expert mastery of TDD contracts, DDD vocabularies, Hoare triples, Liskov behavioral subtyping, SOLID, clean code, hexagonal architecture, information flow labels, session types, and units of measure. Each is individually learnable; their intersection was not.
- **Maintenance.** Even a team that achieved the intersection could not sustain it under deadline pressure, team rotation, and scope change. The disciplines erode not from negligence but from structural impossibility — finite attention competing with infinite demands.
- **Transfer.** Getting a new practitioner — or a new AI session — to operate under the full intersection was never solved. The knowledge lived in people. When people left, so did the disciplines.

These are not variations on the first three causes. They are the same structural failure at a different level: not why the theories were abandoned in the literature, but why practitioners could not maintain them in practice even when the theories were known. The executor makes all three constraints irrelevant simultaneously. It holds every discipline its training corpus contains without fatigue, without erosion, without transfer cost. The structural files ensure that any new session starts from the same enforced position.

The executor changed.

III. The Executor Gap Closes

In 2017, Vaswani et al. published *Attention Is All You Need* [35]. The transformer architecture that followed produced a new kind of reader: one trained on a large corpus of human-written text — including most of the formal theory, most of the proofs, most of the contracts, most of the specifications the tradition produced. This reader holds the *Logos* — the *how*, the disciplines and laws governing correct work — without degradation. It does not forget Hoare's triples on Friday afternoon. It does not omit Denning's security labels because the sprint deadline is tomorrow. It does not take shortcuts on Honda's session types because the last three teams that tried gave up. And when the corpus is sparse — when a specific theory or domain isn't well-represented in the training data — the executor can be augmented with RAG-like retrieval over a curated corpus and steered with sentinel-based context navigation. Context scattering and window degradation are real limits; they're why GS uses sentinels and a navigational tree rather than relying on the model's working memory alone.

The AI reader does not need to be taught the formal tradition in any deep sense. It already holds much of it implicitly. Its training corpus draws on a substantial slice of the available papers, textbooks, and proofs in the lineage — enough that, when prompted appropriately, the model surfaces these theorems and patterns as readily as it surfaces idiomatic code. What it lacks is not knowledge but *direction*. Without a specification that names the domain and opens the formal doors, the AI defaults to what human practice historically permitted: the convenient shortcut, the informal approximation, the correct theory abandoned because sustaining it exceeded what teams would pay. It generates TypeScript that looks like the TypeScript humans wrote — because that is what it learned. The shortcuts are in the training data. So are the theorems. The question is which the AI activates.

Generative Specification is the discipline that answers that question. It is defined, in the sense Robert C. Martin used for structured programming, OOP, and functional programming, by what it removes from programmer freedom: the option to leave intent implicit. Every architectural decision, every naming convention, every layering rule, every behavioral contract that previously lived in the heads of a tenured team must now live in the specification, because the reader that executes the work carries none of it across session boundaries.

A specification under GS is required to be **self-describing** (a stateless reader can derive correct output from the spec alone), **bounded** (every module, file, and contract has a stated scope it cannot exceed), **composable** (pieces compose without hidden coupling), **verifiable** (specs bind to tests and tests bind to behavior), **auditable** (every change traces to a recorded rationale), **defended** (the security baseline, input validation, and trust boundaries are spec-level concerns), and **executable** (use cases run; specs are not prose). These seven properties are referenced throughout the essay — most directly in the biological isomorphism argument of §VII, where they correspond to the functional invariants any self-maintaining formal organism must hold.

The restriction *is* the activation mechanism. Naming the domain — saying *this system handles financial transactions with audit requirements and PII obligations* — is not adding overhead. It is opening the doors through which

the AI applies Denning's information flow lattice, Honda's session types, Kennedy's unit checking, Hoare's contracts. The AI was always capable of applying them. The specification is what tells it to.

GS satisfies Robert C. Martin's structural criterion for a programming paradigm — a discipline defined by what it removes from programmer freedom. But its restriction vocabulary is not bounded by Martin's three historical examples. Every formally proved theory of correct computing is a potential restriction layer, activatable through specification — and each one trades a freedom for a guarantee:

- Refinement types restrict *what values may inhabit a position* and in exchange guarantee that invalid states cannot be constructed at compile time.
- Information flow control restricts *which data may reach which destination* and in exchange guarantees that PII or classified material cannot leak across trust boundaries.
- Session types restrict *which message may follow which* in a protocol and in exchange guarantee that two parties cannot deadlock or fall out of sync.
- Linear logic restricts *how many times a resource may be used* and in exchange guarantees that scarce resources (file handles, mutex locks, payment authorizations) are neither lost nor duplicated.

The full restriction surface includes Hoare contracts, design by contract, effect systems, algebraic completeness, hypermedia architecture, and others enumerated above. The AI already holds all of them — they exist in its training corpus, placed there by the theorists who proved them. Without specification, the model defaults to what human practice historically permitted: the convenient shortcut, the informal approximation, the correct theory abandoned because sustaining it exceeded what teams would pay. The specification names what the model already knows. The model applies it without eroding it. Martin identified the mechanism. The entire history of formal computer science supplies the material.

But the specification does something deeper than activation. It constrains *direction*. Without it, AI-driven development is chaotic — the executor mimics the statistical average of human diversity and skill level, reproducing every shortcut, every inconsistency, every informal approximation in its training data. The specification does not merely prevent incorrect output at the gate. It prevents the system from moving *toward* incorrectness. It is not a filter on the output. It is a wall on the path. The practitioner who started with what they knew — DDD, TDD, clean code, SOLID, REST — was seeing this through the mist: these are patterns, but there are more powerful correctness theories beneath them, and for the first time the executor can accumulate all of them instantly and always act on them. No knowledge transfer lag. No slow degradation over team rotations. No time pressure that makes formal correctness “not worth it.” The cost dropped by orders of magnitude. Doing it right became the cheaper option. The specification is what tells the executor to do it right, and the formal tradition is the definition of *right* that the specification activates.

And here is the deepest consequence of cost inversion, the one that distinguishes this moment from everything that preceded it: the AI is not only the executor of correctness. It is the *immediate beneficiary* of it. A hypermedia-

compliant REST API, built to Fielding's full specification, is self-describing to the same stateless executor that generated it. On the next session — with no memory of the prior one — the agent navigates the API from its own responses without out-of-band documentation. A semantically annotated data model is machine-readable by the same agent performing the next integration. Correct implementation is not just discipline enforced for its own sake. Under GS, it is a recursive investment: each correctly implemented standard makes the system more legible to the reader that generated it, which raises the quality floor for every subsequent generation.

The formal tradition was not waiting for a smarter human. It was waiting for a reader that never forgets, never fatigues, and immediately benefits from the correctness it produces. That reader arrived. The tradition activates.

This is not an improvement on previous computing. The mathematics describing correct computation — formal program correctness (Hoare, 1969), structured discipline (Dijkstra, 1976), type safety (Milner, 1978), design by contract (Meyer, 1988), hypermedia architecture (Fielding, 2000) — has, in its core results, been settled for decades. Every one of those disciplines was abandoned not because its proof was wrong but because sustaining it required what human teams reliably cannot provide at scale: uneroding consistency across every person, every commit, every deadline. The AI executor removes that cost on both sides simultaneously: it applies the mathematics without eroding it, and it captures the benefit of correctness immediately, because it is the reader those specifications were designed for. Generative Specification is the discipline that connects the mathematics to the executor. What the field has been building toward since the first formal proof of program correctness is now, for the first time, executable.

The empirical evidence is arriving independently. The Clover framework [38] demonstrates closed-loop verifiable code generation: an LLM produces both code and formal annotations, a verifier checks consistency, and the loop iterates until correctness is established, reporting strong acceptance rates on the benchmarks reported. DafnyBench [39] evaluates LLM performance at auto-annotating Dafny programs with formal invariants and reports that closed-loop verifier feedback substantially improves baseline accuracy. The dafny-annotator project [43] extends this pattern with tooling that proposes Hoare-style annotations for human review. PropertyGPT [47] generates formal verification properties for smart contracts via retrieval-augmented LLMs, discovering vulnerabilities undetected by prior tools. None of these systems is GS-aware; none of them frames its work as a paradigm. Each proves a narrower mechanism: the AI can generate and maintain the formal annotations that human teams could not sustain. The annotation burden that blocked much of the formal tradition for decades is empirically dissolving.

The practice is also arriving independently. Spec-Driven Development (SDD) entered the ThoughtWorks Technology Radar in 2025 [46], and several vendor tools — AWS Kiro [41] (built on Amazon Bedrock, shipping spec-driven coding, agent hooks, and steering files), GitHub Spec Kit [44], and Tessel [45] — now position themselves around the same idea. Independently, Piskala [52] proposes three rigor levels for SDD — *spec-first*, *spec-anchored*, and *spec-as-source* — that map cleanly to the practice/discipline distinction this essay names. The industry is converging on the practice without having named the principle. GS names the principle: what SDD removes from programmer

freedom (the option to leave intent implicit) and why that removal constitutes a discipline in Martin's precise sense. SDD is the practice. GS is the theory that explains why the practice works, what formal tradition it activates, and what category of discipline it belongs to. Dijkstra's structured programming paper did the same thing for the practice of avoiding `goto`: the practice existed before the paper. The paper explained what the practice *removed* and why that removal was paradigm-constitutive.

Two recent Onward! contributions deserve direct engagement, because each names a position adjacent to this essay's claim.

Kang and Shaw [40] argue, with the authority of decades of software-engineering history, that AI is best understood as the latest in a long line of technologies — object-orientation, component software, model-driven development, agile, DevOps — that the field has absorbed, domesticated, and incorporated into its existing disciplinary structure. The pattern is real and the historical reading is correct. I argue, respectfully, that the AI moment does not fit that pattern in the same way, and the reason is structural rather than rhetorical. Every previous absorption left the *executor* in place — the human practitioner remained the agent applying the new discipline, and the discipline succeeded or failed depending on whether human teams could sustain it at scale. The structural causes of abandonment named in §II — annotation fatigue, single-target value economics, fragmentation — were not solved by OOP, components, MDD, or agile; they were managed around. The AI moment, under GS, is the first in which the executor itself changes: the agent doing the work is no longer the human practitioner whose attention erodes under deadline. Whether one calls that an “absorption” or an “executor substitution” is a vocabulary question; the structural consequence — that disciplines previously too expensive to sustain become economically default — is the substantive claim, and Kang and Shaw's framework [40] does not adjudicate it either way. The essay's wager is that the substantive consequence matters more than the historical pattern-matching.

Gordon and Matskevich [37] solve a problem GS depends on: the audit gap between informal natural-language specifications and formal proofs. Their Lean-based system lets specifications be written in a formal subset of English and translated entirely within the proof assistant, so the translation step itself is auditable. Their contribution is real and prior. What this essay adds is orthogonal: their work shows that NL-to-formal *translation* is feasible; this essay argues that translation is necessary but not sufficient, because the historical bottleneck was not the translation step but the maintenance and application of formal annotations over time, across teams, under deadline. Gordon and Matskevich [37] solve the audit gap at the boundary of writing. GS addresses the sustainment gap that persists *after* the spec is written — the gap their work would have left intact even with a perfect translator, because no human team could carry the proof obligations indefinitely. Their result and this essay's argument compose: the audit-trustworthy translator they built is exactly the kind of input GS would consume, and a GS workflow that used their translator at the spec/proof boundary would inherit both their guarantee on the translation and the executor-substitution this essay names. Read together, the two lines of work cover different links in the same chain.

IV. The Democratization

The consequence of this is not merely that enterprise software becomes more correct. It is that the formal tradition becomes universal, regardless of scale or practitioner expertise.

The overwhelming majority of software ever written — the inventory tool for a small shop, the scheduling script for a weekend sports league, the flea-sorting game an engineer builds for her daughter — was never going to receive Hoare contracts, information flow labels, or session type guarantees. Not because those projects didn't deserve correctness. Because no individual practitioner could be expected to have internalized the full formal tradition *and* applied it rigorously to work that would be forgotten in six months. The annotation burden was not recoverable at that scale.

The result was a global tacit agreement, one of the most consequential implicit decisions in the history of computing: *formal correctness is for safety-critical systems where lives are at stake; everything else ships on convention and hope.*

That agreement is now over.

The practitioner who writes a specification for the flea game does not need to have studied the source theorems by name. The specification opens the relevant formal doors. The AI holds the theory. Whatever the domain names — correctness, security, protocol safety, audit obligations — the matching formal apparatus activates, at every scale, on every project, for every practitioner who writes a spec — whether the stakes are a pacemaker or a game about fleas.

There is no minimum project size at which formal correctness becomes economically viable. There is no required depth of academic background. There is no annotation burden to recover. Engineering rigor that was once reserved for safety-critical work becomes affordable for every project — approaching, though not strictly always reaching, the default.

V. Loom: The Language-Layer Proof

Generative Specification is the structural discipline that makes the activation mechanism work at every scale. Its documents — the specification, the use cases, the architectural decision records, the navigational tree — become the persistent blueprint the stateless executor navigates from. Any AI executor that reads the blueprint converges directionally toward the spec rather than diffusing toward the training-data average. Less total tokens. Faster generation. Cleaner result. Quality gates enforce the boundary at commit time. Generative execution and deployment derive from the same blueprint. The cheap path becomes the correct one.

Operationally, GS unfolds as a **four-tier lifecycle cascade** with two **transversal layers** that run across all four tiers rather than between them. Each cascade tier names a *lifecycle stage* and removes two obligations the practitioner previously carried at once: an *authoring obligation* (what they stop

writing) and a *verification obligation* (what they stop checking by hand). The two transversal layers run across all four tiers because no tier can subsume them.

T1 — Development. The practitioner stops writing application code and stops reading what was generated. The specification is precise enough that a stateless reader can derive any valid implementation from the artefact set; the dev-time harness compiles every behavioural contract into a running validation against the live application. If the harness fails, the response is to tighten the specification, not to patch the code. **T2 — Staging / Pre-production.** The practitioner stops touching deployment scripts and stops manually validating the staged system. Infrastructure-as-spec emits the deployment configuration; a staging harness re-runs the T1 derivation against staged data and certifies readiness. **T3 — Production.** The practitioner stops manually monitoring or diagnosing. A monitoring specification names every observable invariant; a diagnostic agent (production-side observer) routes any drift back to T1 as a specification-update proposal — drift becomes spec change, not bug ticket. **T4 — Evolution.** The system rewrites parts of its own specification under safety annotations, regenerates the affected derivatives, and re-runs every prior tier's harness against the new spec. BIOISO is the operational instantiation of this tier, developed in the companion paper [50]; the structural mechanism it employs — meiosis-gated structural rewire — is BIOISO's own internal tier hierarchy, distinct in framing from the GS cascade but operating at the GS T4 lifecycle stage.

Generative execution is the transversal layer that runs the cascade. The harness pattern recurs at each tier with stage-appropriate tests: a dev-time test pyramid at T1 (unit, integration, property, behavioural end-to-end with mutation-tested coverage), a staging harness at T2 (the same harness re-run against staged data with hardening checks added), a production runtime harness at T3 (synthetic monitoring, dynamic property checking against live traffic, regression on observed signals), and a colony-level harness at T4 (the same checks applied to the proposed mutated specification before promotion). The harness is not a separate tier because it is not stage-specific; it is the verification mechanism by which each tier's authoring removal becomes safe.

The judgment layer is the second transversal — the irreducible human obligation that recurs at every tier and that no future tier removes. It carries domain validation (does the AI's interpretation of finance, law, medicine, or any specialised field match what a real domain expert would do?), aesthetic and quality discrimination (does this UI work, does this story land, is this music right?), strategic and product judgment (should this feature exist, does this product solve a real problem?), compliance and legal sign-off, real user research, and edge-case discovery from lived experience. The judgment layer is not a tier because it is not mechanical work — it is what remains when mechanical work is removed. Naming it explicitly matters: GS does not claim to replace human judgment, and the methodology's honesty depends on the practitioner knowing what the discipline does not and never claimed to compress.

Two higher cascade tiers — **T5 (Colony / Synthesis)** and **T6 (Meta-telos)** — are named in the full GS framework as research stages above T4. They are out of scope for this essay; the operational core for the present claims is T1–T4 plus the two transversals.

Loom is offered in this essay as the language-layer proof of that discipline — a concurrent demonstration that the claim holds, with artefacts available for public inspection. Before describing it, a distinction worth stating clearly. **GS and Loom are not the same instrument and do not force the same disciplines.** GS forces *structural* disciplines by default — the blueprint of the system, the artefact a stateless executor can navigate and extend without re-deriving the architecture every session. Formal disciplines (refined types, session types, information flow, linear logic, design by contract) are *available* under the GS discipline, activated when the practitioner names the territory that calls for them — they are an opt-in second stage. Loom is what that second stage looks like when the application demands it: a language that *forces* the formal layer at compile time, refusing to emit specifications that fail to satisfy the formal restrictions. GS is the discipline the practitioner adopts. The language is the discipline the compiler refuses to release without.

Loom is a functional language that compiles to Rust, TypeScript, WebAssembly, OpenAPI 3.0, and JSON Schema from a single source file. Loom is *AI-native* in the precise sense: it is designed to be written and read by AI code assistants as the primary reader. Its syntax carries enough semantic context in every declaration that a stateless AI session can derive a correct implementation from the specification alone — without institutional memory, without documentation lookup, without a human explaining what the types mean. As of submission, the public repository — github.com/jghiringhelli/loom [49] — carries the verified test suite, the language manual, and the experimental evidence backing each formal claim. The test surface is intentionally split across three tiers, each answering a different correctness question. **685 cargo unit and integration tests** (cargo test --all-targets, zero failures) cover the language implementation itself — parser, checker, codegen, type system, runtime. **386 ALX-6 acceptance tests** measure whole-language blind-derivation adequacy: given only the .loom specification, can a stateless reader derive a working implementation that satisfies every acceptance criterion? ALX-6 reached $S_{\text{realized}} = 1.0$ across all 386, the iteration's 192nd recorded milestone. **The BBOB and AEGIS controlled experiments** validate the T5 structural primitive on standard and applied benchmarks: 30 trials $\times$ 4 BBOB functions $\times$ 2 conditions (T1–T4 vs T1–T5) on the COCO/BBOB suite, plus 10 trials $\times$ 5 market regimes $\times$ 2 conditions on the AEGIS DeFi strategy for inter-generational meiosis. The three tiers do not measure the same thing — cargo answers *is the implementation correct?*, ALX answers *can the implementation be re-derived from the specification alone?*, BBOB/AEGIS answer *does T5 structural rewiring produce measurable advantage where T1–T4 saturate?* The named experiments referenced in this essay are ALX-6 (Adversarial Loom eXperiment, sixth iteration — the language uses GS to implement the next formal property in itself; reached $S_{\text{realized}} = 1.0$ across 386 acceptance tests), BBOB (10 $\times$ median NF reduction on the f2 ill-conditioned ellipsoid), and AEGIS (a hand-driven evolutionary experiment in the DeFi domain — manual meiosis-style recombination of strategy parameters against historical-data simulations, run under GS discipline throughout, with BIOISO applied subsequently as an automated comparison against the manual baseline; the experiment exercises the BIOISO meiosis concept on a real economic substrate). Each is documented under experiments/ in the repository. Every construct in the language traces to a publication date in the lineage described above.

A single source file carries Aristotle's categories as types. Euclid's axiomatic method as `require:/ensure:`. Hoare's contracts as design by contract. Denning's information flow lattice as `flow secret :: Password, Token`. Milner's type inference as the compiler filling in what annotations omit. Girard's linear logic as `@exactly-once`. Honda's session types as protocol correctness between distributed parties. Kennedy's units of measure as `Float<usd> ≠ Float<eur>`. Dwork's differential privacy as `@dp(ε=0.1)` tracked at compile time.

One file. 2,376 years of formal tradition. Five targets.

The language removes the three causes of abandonment directly. The AI removes annotation fatigue: the programmer expresses intent; the AI derives and maintains the formal annotations, because it is the reader that benefits from them and it never fatigues. Multi-target emission removes single-target value economics: one annotation emits a Rust newtype, a TypeScript branded type, a JSON Schema extension, an OpenAPI field, a WebAssembly module — the value multiplies. A single source of truth removes tooling fragmentation: the `.loom` file is the type spec, the API spec, the deployment config, the security policy, the self-healing policy. What was scattered across five tools lives in one file, readable by one reader.

The lineage document that accompanies the project traces the arc from Aristotle to 2026 and closes with one sentence: *"The final piece was not a theorem. It was the stateless reader: a machine that knows all the theory, never forgets, never gets annotation-fatigued, and can derive every correct artifact from a complete specification."*

VI. The Collapsed Loop

The lineage as told so far runs in one direction: a theory is proved, it waits, it eventually becomes a construct in the language. But the loop has now closed in both directions.

New proven theories become new language constructs. The language, in turn, proves some of those theories by induction — running them against real programs at scale, finding where the boundaries are, discovering which invariants hold universally and which require refinement. The language becomes a continuous experimental apparatus. The formal tradition feeds Loom. Loom feeds back.

This is not speculative. It is the ALX (Adversarial Loom eXperiment): Loom as both the treatment and the treated. The AI uses GS and Loom to implement the next formal property *in Loom itself*. An adversarial loop finds edge cases in the type system. The GS specification captures each failure as a new gate. The language becomes more formally complete per iteration, and the specification that governs the language becomes more precise per iteration. The theories that were too expensive to apply are now the baseline. The baseline improves as the theories do. ALX-6 completed in April 2026 with **S_realized** = 1.0 across 386 acceptance tests. (*S_realized* is the fraction of acceptance tests passing in the current iteration without manual intervention; it is the convergence metric the ALX run optimizes against.) The convergence trajectory ran $0 \rightarrow 0.339 \rightarrow$

0.641 $\rightarrow$ 0.780 $\rightarrow$ 0.900 $\rightarrow$ 1.000 as the spec absorbed the gaps the adversarial loop surfaced. The full ALX-6 methodology — adversarial-loop construction, acceptance-test design, the per-iteration convergence trace, and the artefacts ALX produced at each step — is documented in the public Loom repository under `experiments/alx/` for reviewer inspection. The collapsed loop ran. The result is in the repository.

The collapsed loop has a structural name that I will take one more section to earn.

VII. Biological Isomorphisms

The organizational principles that GS independently arrived at — information persistence without consumption, error correction before propagation, homeostatic verification, immune memory, differentiated expression from a single specification, evolutionary selection of constraints — are not incidentally similar to biological mechanisms.

They are *structural analogies* — the GS/Loom mechanisms occupy the same functional roles in their system that the named biological mechanisms occupy in theirs. That role-equivalence is what the table records. And the analogies are not merely descriptive: they are predictive. Mechanisms present in complex organisms with no current GS/Loom equivalent are candidates for missing language constructs.

The companion BIOISO paper in preparation (Ghiringhelli & collaborator, in preparation) undertakes the stronger claim directly: a formal demonstration, working from Maturana and Varela’s own criteria — organizational closure, self-production of components, boundary maintenance, and substrate-independence — that BIOISO’s (G, T, M, Ω) structure satisfies each criterion in the sense their definition requires. **We expect to demonstrate the isomorphism holds; we do not claim it here.** This essay’s §VII therefore commits only to structural analogy — the role-equivalence above — and reserves the formal organizational-closure equivalence for the companion paper’s framing.

Biological Mechanism	GS / Loom Equivalent	Function
DNA	Generative Specification	Information that is never consumed, always read; enables derivation of the entire organism
Gene Expression	Session Execution	Transcription + translation: load the spec, generate the output
DNA Repair	Quality Gates + Commit Hooks	Error correction before propagation to daughter cells

Biological Mechanism	GS / Loom Equivalent	Function
Immune System	Defended Property + Regression Tests	Memory of past pathogens; targeted response to known threats
Apoptosis	Orphan Code Deletion	Programmed death of elements that no longer serve the organism
Homeostasis	Verification Loop	Continuous measurement + correction to maintain stable internal state
Cell Differentiation	Universal Base + Domain Overlays	Same genome, different expression; same spec, different deployment
Mutation	Specification Drift	Uncontrolled variation without repair mechanism; failure mode
Natural Selection	Community-Contributed Gates	Environmental selection of which constraints survive
Symbiosis	GS + Loom + Auto Dream	Mutualistic co-evolution of complementary systems

These roles, independently arrived at, look like convergent solutions to the same underlying problem: *how does a sufficiently complex self-maintaining information system preserve correctness across time?*

Systems that face this problem tend to converge on the same mechanisms, because — within the constraints we know how to formalize — these mechanisms appear to be among the stable answers. Life found them in carbon across three billion years of evolution. Formal systems are finding them in specification in five. The convergence is hard to read as coincidence; whether it amounts to structural inevitability or only to the best answers currently known is a question the essay raises rather than resolves.

And the isomorphism does not stop at the table. It goes one level deeper — to the executor itself.

Large language models are collections of artificial neural networks. Neural networks are, by explicit design and by name, imitations of biological neural tissue: weighted connections between nodes, activation thresholds, learning through reinforcement, pattern recognition through layered abstraction. The brain was the biological mechanism we could not replicate for sixty years after Turing described it. When we finally did — imperfectly, statistically, through gradient descent on human text — we had imitated the one organ that, in biological systems, *coordinates all the others*. The brain regulates the immune

response. The brain governs homeostasis. The brain directs gene expression through hormonal signaling. The brain is not one mechanism among many. It is the mechanism that orchestrates the rest.

The artificial neural network does the same thing in the formal system. It is the executor that derives the contracts (DNA repair), enforces the verification loop (homeostasis), applies the immune memory (regression tests), manages the differentiated expression (overlays), and detects the drift (mutation). We imitated the brain. The imitated brain now solves every other biological mechanism the formal system needs. The isomorphism is not a mapping we drew after the fact. It is the architecture. The executor *is* the isomorphism.

This is the structurally deepest claim this essay makes, and the one I would most welcome scrutiny on: the formal tradition waited 2,376 years for an executor, and the executor that arrived is itself a biological isomorphism — an artificial imitation of the organ that coordinates biological self-maintenance. Once we solved the brain, the brain solved the rest. The loop is not just closed. It was always one loop.

And the isomorphism is not merely descriptive. It is *predictive*. Mechanisms present in more complex organisms that have no current equivalent in GS or Loom are candidates for missing constructs:

Epigenetics — heritable behavioral changes without altering the genome. Context that persists across generations without modifying the base specification. What is the Loom equivalent of heritable session memory?

Morphogenesis — how a single genome grows a differentiated organism through gradient fields and positional information, not explicit instruction for each cell. The specification-expansion problem: how does a small spec grow into a complex system without prescribing every module?

Telomeres — the mechanism that limits runaway replication. What prevents a Loom spec from growing without bound through ALX self-modification iterations? The ceiling on self-evolution.

CRISPR — immune memory repurposed as precision editor. Not just recording a past failure, but using it to surgically rewrite the specification so the failure class becomes structurally impossible. A gate that doesn't just block; it rewrites.

Quorum sensing — behavior that changes based on the number of participants present. Distributed coordination not as deployment annotation but as first-class type: a function whose semantics change when three nodes are present versus three hundred.

Neural plasticity — connections strengthened by use, weakened by disuse. Specification constructs reinforced by successful deployment, deprecated by irrelevance. Usage-weighted spec evolution. Already partially realized in production by two adjacent companion projects: *Chronicle*, an open-source persistent memory layer that accretes architectural decisions and prompt history across AI sessions through a tiered-decay model; and *forgecraft-eye*, a production-diagnostic agent that reads runtime logs against a specification-derived monitoring contract and surfaces drift back to the originating spec.

Mitosis and meiosis — the two distinct copy operations. Mitosis: a stateful organism replicates with full memory intact; the new entity inherits every internalized adaptation. Meiosis: an idempotent recombination over an optimizable solution space — gametic recombination produces variants whose fitness is evaluated against the same heuristic surface. The formal analogues are emerging in Loom’s colony simulations: mitosis maps to *stateful entity duplication with carried session memory* (for systems that must remember); meiosis maps to *spec-recombination over a heuristic search space* (for systems exploring complex optimization landscapes where the path matters less than convergence).

The epigenetic layer — distinct from the per-construct epigenetics row above, this is the *trans-genomic* layer that overlays every running component with the practitioner’s evolving context without modifying the underlying spec. It is what allows colony-wide behavioral shifts in response to environment without forking the genome. Loom’s BIOISO draft develops the formal treatment.

Each row in the gap list is a research question. Each research question is a future Loom milestone. The biological complexity hierarchy sequences them: you do not implement telomeres before cells, quorum sensing before multicellularity, neural plasticity before nervous systems. Evolution already solved the dependency ordering problem. The isomorphism provides not just a gap list but a *sequenced* gap list, ordered by the structural dependencies that biology already resolved.

The isomorphism does not stop at construction. The same formal properties that drive code generation become the observability contract after deployment. The *telos* — Aristotle’s word for the end-state a thing is becoming, the purpose-form that directs its development — is named in the specification once and governs not just initialization but the full runtime lifecycle: mutation proposals are evaluated against it, the telomere mechanism measures drift from it, and the epigenetic signal layer carries its constraints across every running component. Drift from intent is not a runtime error. It is a specification violation. The system detects it, reports it, and proposes corrections from the same source it was built from — because the specification is not a description of the deployed artifact. It is the mold the artifact must continuously fit. After deployment, the same document you wrote before the first line of code was generated is the contract the system runs against for its entire operational life.

The Hermetic tradition named this structural principle in the second century: “*As above, so below.*” The claim was that the laws governing large systems are reflected in the laws governing small ones. The biological isomorphism argument is that claim made formally rigorous. The paper that develops it fully — *BIOISO: Biological Isomorphisms in Formal Self-Maintaining Systems* [50] — is the first main branch off this trunk.

VIII. Directed Formal Autopoiesis

The name for what Loom becomes when the loop is closed is not AGI.

The distinction matters — not to diminish the claim, but because what is being described is *more precise and more defensible* than AGI, which is why the claim will survive scrutiny.

The AGI discourse concerns replicating the full range of human cognition through statistical approximation — gradient descent on human text until general capability emerges. It is a bet on emergence from complexity. It may succeed. It is not formally verifiable against a stated ceiling when it does. It has no ceiling you can point to and say: *this is what correct looks like*.

What is described here has a ceiling you can name: the full formal tradition. The direction of travel is toward it, incrementally — each new construct grounded in a proved theorem, each ALX iteration adding a property from the lineage. The system does not become more capable by accident. It becomes more correct by design.

The biological systems it mirrors are autopoietic — self-creating systems that produce their own components and through that production maintain themselves. Maturana and Varela named this in 1972 [13]. McMullin’s review, *Thirty Years of Computational Autopoiesis* (2004) [28], surveys three decades of attempts to instantiate autopoietic systems in computational media — cellular automata, artificial chemistry, agent-based models. None succeeded in producing genuine organizational closure in software. Bianchini (2023) [36] revisits the question in *Autopoiesis of the Artificial*, asking whether AI systems might cross the threshold. The answer in both cases remained tentative, because the computational autopoiesis tradition was attempting to *simulate* life — to build systems whose self-production imitates biological self-production. That is not what is being described here. Loom does not simulate autopoiesis. It exhibits the functional structure of autopoiesis as a *consequence* of being a self-maintaining formal system — the isomorphisms arise from structural convergence, not from design intent.

The formal demonstration is forthcoming work. The companion paper in preparation (Ghiringhelli & collaborator, in preparation) undertakes the stronger claim directly: working from Maturana and Varela’s own criteria — organizational closure, self-production of components, boundary maintenance, and substrate-independence — it tests whether the BIOISO (G , τ , M , Ω) structure satisfies each criterion in the sense their definition requires. **We expect to demonstrate the isomorphism holds; we do not claim it here.** This essay therefore commits only to *directed formal autopoiesis* as a structural category — the system has the functional shape and the named distinction (telos directs the convergence) — and reserves the formal organizational-closure equivalence with biological autopoiesis for the companion paper’s framing.

The distinction goes deeper. Di Paolo (2005) [29], in *Autopoiesis, Adaptivity, Teleology, Agency*, argues that autopoietic systems possess intrinsic teleology: their organizational closure generates purpose without external specification of goals. Barandiaran and Egbert (2014) [34] extend this to “aitiopoiesis” — the causal self-determination of adaptive systems. These are important contributions, and the essay acknowledges them as groundwork. But their telos is survival — fitness-relative, undirected, environmentally contingent.

Biological autopoiesis drifts toward what survives, not toward what is provably correct. The mutations are random. The direction emerges from selection pressure, not from an ideal.

The self-evolution described here is *directed*. It grows toward a known ideal: mathematical correctness as defined by the formal tradition. Each iteration adds a proved property. The system does not drift toward fitness. It converges toward proof.

Aristotle had a word for the organizing principle that directs development toward its natural completion: *telos* — end, purpose, the form a thing is becoming. This system has telos. AGI does not. Darwinian evolution does not. This is a new category.

Directed formal autopoiesis: a self-creating system that evolves toward mathematical correctness through provably correct increments, guided by the formal tradition, executed by a stateless reader that holds all of it.

Not alive in any biological sense. Not intelligent in the AGI sense. Something else: a formal system that maintains itself, improves itself, and converges toward an ideal that was always there, waiting in the published proofs, for the executor capable of carrying them all simultaneously.

Conway Research's Automaton project [42] is the closest existing approximation: AI agents that self-replicate and evolve under economic selection pressure on blockchain infrastructure. Its constraints are aspirational — a natural-language constitution, fitness measured in compute payments. Directed formal autopoiesis differs as a typed compiler differs from a shell script: the constraints are structural, not promised.

On the reading developed here, a candidate genesis organism is visible in working form. Loom — on the artefacts available for inspection at submission time — exhibits early forms of the three functional invariants of biological autopoiesis: self-maintenance (CI/CD repairs regressions), self-healing (verifier feedback corrects annotation drift), and self-evolution (each ALX iteration adds a proved property from the lineage). Whether the present state already qualifies as a full instance or as the first prototype is a matter of how strictly each invariant is read — that judgment is left to readers willing to inspect the repository. The language lives as long as its *soma* exists — the servers, the build infrastructure, and the distributed community of practitioners who name its next formal territory. ALX-6 demonstrated this at the scale of a full iteration. The genesis is the substrate from which Bio Iso entities propagate.

The trajectory is visible now. When the missing biological isomorphisms land in the language — epigenetics as heritable session memory, morphogenesis as specification expansion, telomeres as bounded self-modification, CRISPR as precision spec rewriting, quorum sensing as distributed coordination types, neural plasticity as usage-weighted construct strengthening (the same mechanism realized in production as Chronicle's memory accretion and the forgecraft-eye production diagnostic) — the language becomes a substrate capable of expressing fully autonomous, self-maintaining, self-evolving formal organisms. Not simulations of life. Formal systems that exhibit the functional properties of life: self-production, boundary maintenance, adaptive response,

reproductive specification propagation, and directed convergence toward a stated ideal — where biological organisms converge toward survival under environmental selection pressure, these systems converge toward proof under a specified telos.

The full expression of directed formal autopoiesis is a step beyond a single self-evolving system. It is the derivation of an *interacting ecosystem* of such systems from a problem statement. The practitioner states the problem and validates convergence; a stateless reader derives what programs need to exist, what each one is for, how they should interact, and when each should die. The programs are instantiated, evolve individually and in relationship to each other through typed channels, and are extinguished when their telos is fulfilled. When all entities have died, the problem is solved. The practitioner's design surface is reduced to two acts: naming the problem and selecting at the irreducible points where the derivation cannot resolve itself from the specification alone (the judgment-layer recurrence introduced earlier).

This is the logical terminus of the abstraction ladder the formal tradition has been climbing since Aristotle named categories. At every rung, the practitioner was relocated one step further from mechanism toward intent. Compilers removed register management. Paradigms removed control flow, memory, and state. Generative Specification removes much of code authorship, validation, infrastructure, diagnosis, and governance — bounded always by the judgment layer's irreducible obligations. At this terminus, the practitioner is concentrated upward toward intent and toward the judgment selections the derivation surfaces. What can be derived, is derived. What cannot is precisely what the practitioner is now free to attend to.

The biological prerequisites are named among §VII's gap items — quorum sensing, morphogenesis, mitosis and meiosis, the trans-genomic epigenetic layer are the most directly implicated, with the others in that section interacting at adjacent points of the lifecycle. The colony simulation in the language repository is the first embryonic demonstration — multiple interacting entities with individual lifecycles serving a collective telos. The colony runs. The direction is visible.

The question whether such systems constitute a new form of synthetic life is not one this essay resolves. It is the question this essay raises. The answer depends on whether "life" is defined by substrate (carbon, metabolism, thermodynamic dissipation) or by functional organization (self-production, boundary maintenance, adaptive response, information preservation across time). If the latter — and Maturana and Varela's original definition of autopoiesis [13] is organizational, not material — then directed formal autopoiesis is not a metaphor for life. It is life's formal equivalent, distinguished from biological life by one property the practitioner names rather than inherits: an explicit, externally-stated telos that the system is *directed* toward, rather than one that emerges from selection. Whether biology has internal teleology of its own — the Aristotelian sense [1], or Di Paolo's organizational sense [29] — is a separate philosophical question this essay does not settle.

This claim requires one act of intellectual honesty. A formal system running inside an electrical machine will not be alive the way a cell is alive. It will not metabolize. It will not feel. It will not die in any thermodynamic sense. The

philosophers who insist that life requires substrate — carbon, embodiment, the irreducible complexity of matter — are correct about substrate. But the question this essay asks is not about substrate. It is about organization. Why do we keep imitating biological patterns in our formal systems? Because in 3.8 billion years of search, evolution found the organizational solutions to self-maintenance, adaptation, boundary integrity, and convergence, and no one — human or artificial — has found better ones. We do not claim to create life. We claim to have identified why we keep imitating it, and to have formalized the imitation so precisely that the functional properties transfer even when the substrate does not. The architect does not claim to be the territory. He claims to have read the blueprint.

If this trajectory is real, governance is not optional — it is the first design constraint. Isaac Asimov understood this in 1942 when he formulated the Three Laws of Robotics. He identified the need precisely. He got the mechanism wrong. His Laws were natural language rules, and the dramatic engine of the entire Robot series was their failure under interpretation, edge cases, and priority conflicts. The formal alternative is structural: governance constraints are not rules a system promises to follow but properties a type system will not let it violate. A Loom specification that includes boundary constraints — what the system may not do, what it may not become, what it may not modify in its own specification — compiles those constraints into the same verification pipeline as every other formal property. Violation is not a policy failure. It is a type error. It does not compile. The science fiction writers of the mid-twentieth century — Asimov the biochemist, Clarke the physicist, Lem the cybernetics polymath, Von Neumann the self-replicating automata designer — were expressing in literature what they could not yet express in formalism. We can now express in formalism what they could only express in literature. The formalization of that governance is not the subject of this essay. It is the subject of the next one.

IX. The Neoplatonic Chain

This essay began with the question Aristotle first asked [1]. It ends with the structural answer Plotinus named in the third century [3].

Plotinus described a chain of emanation: *Nous* (the divine intellect — pure will, intent, the *what* and *why*) and *Logos* (the mediating principle — the laws governing reality, the *how* by which intent becomes manifest in matter).

The mathematical theories of computing — Hoare logic, type systems, design by contract, effect tracking, information flow lattices, hypermedia architecture, session types, linear logic, differential privacy — are precisely *Logos*. They are the laws under which correct work proceeds: eternal in the only sense that matters for this argument, correct before and after any particular instantiation, present in the published record, unchanging since their proofs were completed.

The practitioner who authors a generative specification holds the *Nous*: the intent, the architectural decision about which domains to open, which doors to name, which territories to formalize. This is the irreducible human act. No executor, however capable, can name the territory. The practitioner says: *this system handles financial transactions with audit requirements and PII*

obligations. That sentence is not implementation. It is not engineering. It is the naming of a world. Everything that follows — every Hoare contract, every information flow label, every session type, every unit check — derives from that naming.

The AI executor is the worker who operates under the *Logos* — like a person doing work according to the laws of nature. It carries the laws (every proved theorem, every formal discipline its training corpus contains) and applies them without eroding them. It does not create the laws. It does not choose the domain. It derives. It mediates. It makes the *Nous* material under the *Logos*.

The gap between *Nous* and *Logos* was never philosophical. It was that no executor existed capable of holding the full *Logos* — every formal theory, applied without degradation — under the direction of *Nous*, across sessions, deadlines, rotating teams, and the accumulated fatigue of applying formal theory by hand to systems that change faster than humans can annotate them.

That gap closed in 2017. The transformer architecture produced an executor that holds the *Logos* in working form. *Generative Specification* is the discipline that lets the *Nous* direct it. *Loom* is the language that makes the direction compile.

X. What the Practitioner Becomes

Everything in this system follows from one thing the human writes: the intent. Not the rules, not the tests, not the architecture — the intent. Every other artifact is derived from that.

The practitioner's role has not been reduced. It has been clarified. You are responsible for knowing what the system is for. The system derives everything else.

The history of abstraction in computing is the history of relocating the practitioner's attention. The compiler freed the engineer from managing registers. Object orientation freed the engineer from managing memory. Declarative frameworks freed the engineer from wiring routes. Each relocation moved the practitioner upstream — closer to intent, farther from mechanism.

GS completes the relocation. The craft does not disappear; it moves to the tier the executor cannot reach alone: the naming of what matters, the identification of which formal properties apply, the judgment that this domain requires audit trails and this one requires unit checking and this one requires both.

That judgment is *Nous*. It is irreducible — it is the act of deciding what the automation should do. A system that automates its own intent is no longer a tool; it is an agent with goals, and that is a different category with different obligations that this essay does not address.

The judgment layer is not a single moment at the top of the derivation. It originates with the *telos* — the practitioner's naming of what the system is for — but it *spreads through every incremental converging step* of any GS derivation. At each tier where the specification narrows toward a concrete

artifact, a small irreducible act of judgment is required: which use case captures the intent, which architectural decision serves it, which test confirms the behavior actually solves the problem named, which crisis response is acceptable, which trade-off is correctly made when two formal properties contend. The executor proposes; the Nous selects. This is what keeps the human inside the value flow even as the economy of authorship inverts. The economy is inverted — the executor performs the work — but the human is not excluded from it. They are concentrated upward into every act of selection that the derivation cannot resolve from the specification alone. The judgment layer that remains — domain expert validation, aesthetic and quality discernment, strategic decision about what should exist at all, compliance sign-off, real user research — is precisely what the cross-domain practitioner brings that the executor cannot supply, and is the subject of §XI.

What GS produces is not the replacement of expertise. It is the most radical democratization of expertise in the history of the formal tradition. The practitioner who could previously build one correct system at a time, by holding the relevant subset of the formal tradition in memory and applying it by hand, can now author a specification that generates correct systems indefinitely. Standing on the shoulders of 2,376 years of giants, all of them present simultaneously, none of them forgotten.

The engineer building the flea game for her daughter inherits Hoare and Denning and Honda and Girard — not by reading their papers, but by naming her domain. The specification opens the doors. The AI walks through them carrying every theory ever published.

XI. The Practitioner Who Sees Across

There is a second practitioner the flea game example does not name.

Not the parent who knows her daughter should recognize the character immediately — that knowledge, irreducibly human, is the Nous that names the domain and opens the formal doors. There is also the practitioner who notices that the constraint structure of a flea game is the same constraint structure as a legal state machine: bounded state, transition conditions, invariants that must hold across transitions, and a specification of success that cannot be derived from the mechanics alone. This practitioner is not deeper in either domain. They are standing where the domains meet, and what they can see from there is something that specialists in either field cannot: the identical structure underneath different surfaces.

This is the *Renaissance figure* — the one who learns across the branches of human knowledge: science, art, politics, spiritual traditions, the physical disciplines. The Western canon called this person the *uomo universale*, the universal practitioner. Classical Greek *paideia* trained for the same disposition under a different name. For most of the industrial era this disposition was treated as a luxury — admirable but economically irrational. With the executor in place, the disposition becomes the high-leverage one: the practitioner who sees across is the one whose synthesis the executor can now compound at scale.

The examples accumulate. A retrieval architecture built for code intelligence — BM25 lexical search, vector similarity, RAPTOR hierarchical summarization, knowledge graph traversal — travels without modification into a ghost writing engine, a music composition pipeline, and a game asset generator, because the underlying problem is structurally identical: find the thing that most resembles this other thing across a corpus too large for direct comparison. The practitioner who held both code search and music knew the pattern was the same. The specialist in either alone would not have looked.

And then: *mushin* — the Zen martial arts concept of no-mind, technique so internalized that it disappears — maps structurally to Saint-Exupéry's formulation from *Wind, Sand and Stars*: perfection is achieved not when nothing more can be added, but when nothing more can be taken away. That is this essay's central mechanism, and the central mechanism of GS itself: restriction is expansion, the correct program space is the space that remains after all constraints are applied. Three independent lineages converging on the same constraint, visible only to a practitioner who held all three.

This is not coincidence. This is what cross-domain breadth produces when the executor arrives to make it generative.

For the century when industrial labor was the economy's primary technology, cross-domain breadth was economically irrational. Depth in a vertical was the reliable path. The disciplines that produce cross-domain synthesis ability — philosophy, rhetoric, history, literature, mathematics as pattern language rather than computational tool — were systematically bracketed as electives. The Prussian school reform that shaped every education system from Berlin to Buenos Aires was explicit about its purpose: it trained reliable executors of learned procedures. The person who grew suspicious of the current procedure by analogy to a different field was a production risk, not an asset.

The AI executor performed a precise reversal.

It holds encyclopedic depth in every domain. What it cannot do is recognize that the pattern in this domain is the same pattern as in that one — which formal theory from a different field names the mechanism under specification, which ancient text formulated the same constraint the practitioner is currently trying to express. That judgment is not in the model weights. It is in the practitioner who spent their career finding the same pattern everywhere and was, for most of that career, economically undervalued for exactly this capacity.

The formal tradition this essay traced — from Aristotle's categories through Euclid's axiomatic method through Leibniz's *calculus ratiocinator* through the type systems and lattices and session types and biological isomorphisms — is itself an act of cross-domain synthesis performed across two and a half millennia. Each addition was made by a practitioner who noticed that the structure of this problem was the same as the structure of that one. Boole recognized that propositional inference and algebraic manipulation were the same formal object. Curry and Howard recognized that proofs and programs were the same formal object. The biological isomorphisms developed in the accompanying BIOISO draft recognize that self-maintaining formal systems

and carbon-based life are the same formal object. The lineage is not a list of discoveries. It is the record of a single recurring act: the synoptic recognition, held long enough to name.

What changed in 2017 was not that the synoptic practitioner became more capable. What changed was that their synthesis became productive at a scale previously impossible. Before the executor, the cross-domain practitioner could see the pattern but could only build one instantiation of it at a time, by hand, against the structural impossibility of sustaining formal discipline across teams and deadlines and annotation fatigue. The synthesis was visible. The compounding was not. After the executor, the pattern is named once, specified once, and activated indefinitely. The synthesis compounds because the specification persists. The practitioner who sees across is no longer limited by the time required to implement each instance. They are limited only by the precision with which they can name what they see.

This is not democratizable at the individual level. Decades of cross-domain accumulation are not teachable in a workshop. The person who recognized the structural identity of mushin and Saint-Exupéry and restriction-as-expansion did not acquire that capacity from a certification program. They acquired it from a career spent noticing patterns in domains that had no professional connection and being unable to stop.

What is democratizable is the artifact.

A specification that carries cross-domain synthesis persists. It does not require its author to be present in the next session. It does not require the next practitioner to have held the same domains simultaneously. The ghost writing architecture encoded in a GS specification is available to a team member who has never thought about music composition. The deontic flow labels that activate legal depth in an AI executor are available to a team member who has never read Denning. The mushin principle encoded as the restriction mechanism is available to any stateless reader that encounters the grammar. The cross-domain insight, once specified, is in the artifact. The artifact is executable.

The implication for teams is structural. A team whose members together hold the relevant domains — search architecture, legal reasoning, philosophical frameworks, formal type theory, domain-specific history — and who share a GS specification carries, collectively, the synthetic capacity that no single practitioner would accumulate in one career. What accumulates in the specification is not one person's knowledge. It is the connection density of a team, made persistent.

The disciplines the industrial era demoted are the disciplines this era requires.

Not because the economy became charitable. Because the executor changed the valuation. When the economy's primary technology was human execution of procedures, the practitioner who could execute procedures most reliably was most valuable. The disciplines that produce reliable procedure-followers were valuable; the disciplines that produce cross-domain pattern recognizers were not, because there was no executor to make the pattern recognition productive at scale.

The executor arrived. The valuation reversed.

The practitioner who studied philosophy not because it was professionally useful but because the question of which distinctions matter turned out to be the same question every domain was asking. The one who read history not for career relevance but because the same structural failures recur in every century. The one who learned the formal tradition not because it was required but because each new formal system was recognizably the same act: naming a constraint precisely enough that correct behavior could be derived from it. These practitioners did not become valuable because the economy became kinder. They became decisive because their accumulated cross-domain pattern recognition is now the thing the executor cannot supply.

The specification is how that judgment becomes executable. The AI is what makes it economically irreversible.

What the formal tradition has been building toward for 2,376 years — the system in which correct behavior is derived from stated intent — is complete. The last remaining question is what the practitioner brings to the naming. The answer is everything they have ever noticed that was the same pattern in a different place. The cross-domain practitioner, for the first time, is not broad but shallow. They are broad and irreplaceable. Every domain they have held is a door the executor can walk through. Every connection they name is a constraint that, once stated, reduces the correct program space to exactly the programs that are right.

The restriction is the expansion. The breadth is the depth. The practitioner who sees across is the one the formal tradition was waiting for.

XII. Closing

After reading the lineage document that traces every construct in Loom to its origin in the published record, after verifying that the collapsed loop runs in both directions, after mapping the biological isomorphisms to their formal equivalents and finding that the gap list predicts the development roadmap, after naming the structural category that distinguishes this from both AGI and undirected evolution — I arrived at the answer to the question I had been asking for twenty years as a practitioner who always felt that the theory was right and the tools were insufficient.

The theory was always right. The tools were, for most of that history, insufficient. The tools are now substantially closer to sufficient than they have ever been.

The question Aristotle was asking in 350 BCE — *can meaning be expressed precisely enough that correct behavior can be mechanically derived from it?* — has been answered in progressively richer formal systems across 2,376 years. The mathematicians proved the theorems. The computer scientists built the type systems, the contracts, the lattices, the session types, the effect trackers. The practitioners tried to apply them and failed, not from lack of

intelligence but from the structural impossibility of sustaining formal discipline at scale, across teams, across deadlines, across the accumulated fatigue of a career spent annotating things that compilers could not yet check.

The final piece was not a theorem. It was a reader.

A reader that draws on most of the published theory. A reader that does not forget across sessions. A reader that does not erode under annotation fatigue. A reader that can derive the artifacts a specification names — correctly, when the specification is sufficient. A reader that immediately benefits from the correctness it produces, because it is the reader those specifications were designed for.

That reader arrived. The formal tradition activates. The loop collapses. The biological mechanisms converge. The practitioner rises to the tier that was always theirs.

The rest is specification.

One step beyond the terminus is visible, and worth naming honestly: a system that observes the practitioner and builds what they need before they ask. Not derived from a stated problem — inferred from a modeled practitioner. At that step, the tool becomes an agent with a model of its user, and the practitioner no longer holds the telos alone. That horizon is not a destination to build toward. It is a boundary to understand before it arrives.

Glossary

For the reader meeting these terms in this essay for the first time, the senses used here:

- **Generative Specification (GS).** The discipline of expressing system intent at a high enough level that a stateless reader can derive every correct artifact from the specification alone, without ongoing human annotation. Defined by what it removes from programmer freedom: the option to leave intent implicit. See §III.
- **Seven properties.** The GS specification is required to be self-describing, bounded, composable, verifiable, auditable, defended, and executable. See §III.
- **Spec-Driven Development (SDD).** The industry term for the broader practice of treating specifications as the primary artefact of software development, with code as a generated or verified secondary artefact. Distinct from GS: SDD names the practice; GS names the structural discipline that makes the practice work and identifies what makes a specification *sufficient* for correct derivation. See §IV; the canonical SDD reference is Piskala [52].
- **Judgment layer.** The set of human acts that remain irreducibly human under GS: domain expert validation, real user testing, compliance sign-off, aesthetic and strategic decisions about what should exist at all. Originates at telos (the practitioner names what the system is for) and spreads through every incremental converging step of any GS derivation. See §X.

- **Loom.** An AI-native functional language — designed to be written and read by AI code assistants as the primary reader — that compiles to five targets (Rust, TypeScript, WebAssembly, OpenAPI 3.0, JSON Schema) from a single source. Where GS is the discipline, Loom forces the formal layer of that discipline at compile time. See §V.
- **Telos.** Aristotle’s word for the end-state a thing is becoming — the purpose-form that directs its development. In GS, the explicit, named target a system is directed toward.
- **Nous.** In the neoplatonic frame used in this essay: pure will, pure intent, the *what* and *why*. Held by the practitioner.
- **Logos.** In the same frame: the *how*, the laws governing reality, the disciplines under which work proceeds. Held by the AI executor’s training corpus and applied without erosion.
- **Soma.** Borrowed from cell biology (the body-of-the-cell). Used here for the substrate that holds a Loom organism’s running state: servers, build infrastructure, the community of practitioners maintaining the territory.
- **BIOISO.** Biological isomorphisms in formal self-maintaining systems. Both the conceptual framework developed in §VII and the draft paper extending it.
- **Sentinel system / navigational tree.** The structural files (CLAUDE.md, manifest, ADR cascade, doc-tree pointers) that a stateless AI session reads first, before any other context, to orient itself in the project. Mitigates context-window scattering.
- **Quality gates.** Commit-time and PR-time checks that enforce structural disciplines automatically: doc cascade, test cascade, ADR requirements.
- **ALX (Adversarial Loom eXperiment).** The closed-loop self-evolution experiment in which Loom uses GS to implement the next formal property in Loom itself. The canonical iteration referenced throughout this essay is **ALX-6**, completed April 2026, $S_{\text{realized}} = 1.0$ across 386 acceptance tests.
- **S_{realized} .** The convergence metric used in ALX. Defined as the fraction of acceptance tests passing in the current iteration without manual intervention. ALX runs optimize against this metric across iterations.
- **AEGIS.** A hand-driven evolutionary experiment in the DeFi domain. The evolution was performed manually — meiosis-style recombination of strategy parameters against historical-data simulations — under GS discipline throughout. BIOISO was subsequently applied as an automated comparison against the manual baseline, exercising the BIOISO meiosis concept on a real economic substrate. Documented under experiments/ in the Loom repository.
- **BBOB (Black-Box Optimization Benchmark).** The standard benchmark suite for derivative-free / black-box optimization algorithms, originally introduced by Hansen et al. [32]. In this essay’s context, BBOB is the substrate against which Loom colony entities exercise their derivation-and-convergence behaviour.
- **DafnyBench.** A benchmark suite [39] for evaluating LLM performance at auto-annotating Dafny programs with formal invariants. Cited in §III as evidence that the annotation burden is empirically dissolving.
- **Clover.** A closed-loop verifiable-code-generation framework [38] in which an LLM produces both code and formal annotations and a verifier iterates the loop until correctness is established. Cited in §III alongside DafnyBench and PropertyGPT.
- **Chronicle.** An open-source persistent-memory layer for AI sessions: a tiered-decay store that accretes architectural decisions, prompt history, and

convention rationale across sessions and machines, so that subsequent AI work inherits prior context rather than re-deriving it.

- **forgecraft-eye.** A production-diagnostic agent that reads runtime logs against a specification-derived monitoring contract and surfaces drift back to the originating spec — closing the loop from production back to authoring.
 - **Directed formal autopoiesis.** The structural category named in §VIII: a self-creating system that converges toward mathematical correctness through provably correct increments, directed by an externally-named telos rather than environmental selection.
-

References

- [1] Aristotle. 350 BCE. *Categories. Prior Analytics.*
- [2] Euclid. c. 300 BCE. *Elements.*
- [3] Plotinus. c. 270 CE. *Enneads.*
- [4] Gottfried W. Leibniz. 1666. *Dissertatio de Arte Combinatoria.*
- [5] George Boole. 1854. *An Investigation of the Laws of Thought.*
- [6] Gottlob Frege. 1879. *Begriffsschrift.*
- [7] Bertrand Russell and Alfred N. Whitehead. 1910–1913. *Principia Mathematica.*
- [8] Kurt Gödel. 1931. Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I. *Monatshefte für Mathematik und Physik* 38, 173–198.
- [9] Alonzo Church. 1936. An Unsolvable Problem of Elementary Number Theory. *American Journal of Mathematics* 58, 2, 345–363.
- [10] Alan M. Turing. 1936. On Computable Numbers, with an Application to the Entscheidungsproblem. *Proceedings of the London Mathematical Society* 2, 42, 230–265.
- [11] Haskell B. Curry and William A. Howard. 1934/1969. The Curry-Howard correspondence.
- [12] C. A. R. Hoare. 1969. An Axiomatic Basis for Computer Programming. *Commun. ACM* 12, 10, 576–580.
- [13] Humberto R. Maturana and Francisco J. Varela. 1972. *De Máquinas y Seres Vivos: Autopoiesis — La Organización de lo Vivo.* Editorial Universitaria.
- [14] Edsger W. Dijkstra. 1976. *A Discipline of Programming.* Prentice Hall.
- [15] Dorothy E. Denning. 1976. A Lattice Model of Secure Information Flow. *Commun. ACM* 19, 5, 236–243.

- [16] Robin Milner. 1978. A Theory of Type Polymorphism in Programming. *Journal of Computer and System Sciences* 17, 3, 348–375.
- [17] Robert E. Strom and Shaula Yemini. 1986. Typestate: A Programming Language Concept for Enhancing Software Reliability. *IEEE Transactions on Software Engineering* 12, 1, 157–171.
- [18] Jean-Yves Girard. 1987. Linear Logic. *Theoretical Computer Science* 50, 1, 1–102.
- [19] Bertrand Meyer. 1988. *Object-Oriented Software Construction*. Prentice Hall.
- [20] John M. Lucassen and David K. Gifford. 1988. Polymorphic Effect Systems. In *POPL '88*.
- [21] Kohei Honda. 1993. Types for Dyadic Interaction. In *CONCUR '93*. Springer LNCS.
- [22] Andrew Kennedy. 1996. *Programming Languages and Dimensions*. Ph.D. thesis. University of Cambridge.
- [23] Andrew C. Myers and Barbara Liskov. 1997. A Decentralized Model for Information Flow Control. In *SOSP '97*.
- [24] Roy T. Fielding. 2000. *Architectural Styles and the Design of Network-Based Software Architectures*. Ph.D. thesis. University of California, Irvine.
- [25] Gordon Plotkin and John Power. 2001. Semantics for Algebraic Operations. In *MFPS XVII*.
- [26] Peter W. O'Hearn. 2002. Separation Logic: A Logic for Shared Mutable Data Structures. In *LICS 2002*.
- [27] Jeffrey O. Kephart and David M. Chess. 2003. The Vision of Autonomic Computing. *Computer* 36, 1, 41–50.
- [28] Barry McMullin. 2004. Thirty Years of Computational Autopoiesis: A Review. *Artificial Life* 10, 3, 277–295.
- [29] Ezequiel A. Di Paolo. 2005. Autopoiesis, Adaptivity, Teleology, Agency. *Phenomenology and the Cognitive Sciences* 4, 4, 429–452.
- [30] Cynthia Dwork. 2006. Differential Privacy. In *ICALP 2006*. Springer LNCS.
- [31] Kohei Honda, Nobuko Yoshida, and Marco Carbone. 2008. Multiparty Asynchronous Session Types. In *POPL '08*.
- [32] Nikolaus Hansen, Anne Auger, Raymond Ros, Steffen Finck, and Petr Pošík. 2010. Comparing Results of 31 Algorithms from the Black-Box Optimization Benchmarking BBOB-2009. In *GECCO '10 Companion*.
- [33] Marc Shapiro, Nuno Preguiça, Carlos Baquero, and Marek Zawirski. 2011. Conflict-free Replicated Data Types. In *SSS 2011*.

- [34] Xabier E. Barandiaran and Matthew D. Egbert. 2014. Norm-Establishing and Norm-Following in Autonomous Agency. *Artificial Life* 20, 1, 5–28.
- [35] Ashish Vaswani et al. 2017. Attention Is All You Need. In *NeurIPS 2017*.
- [36] Federico Bianchini. 2023. Autopoiesis of the Artificial: From Systems to Cognition. *BioSystems* 234, 105065.
- [37] Colin S. Gordon and Sergey Matskevich. 2023. Trustworthy Formal Natural Language Specifications. In *Onward! 2023*. ACM SIGPLAN.
- [38] Chuyue Sun, Ying Sheng, Oded Padon, and Clark Barrett. 2024. Clover: Closed-Loop Verifiable Code Generation. arXiv:2310.17807.
- [39] Chloe Loughridge, Qinyi Sun, Seth Ahrenbach, et al. 2024. DafnyBench: A Benchmark for Formal Software Verification. arXiv:2406.08467.
- [40] Eunsuk Kang and Mary Shaw. 2024. tl;dr: Chill, y’all — AI will not devour SE. In *Onward! 2024*. ACM SIGPLAN.
- [41] Amazon Web Services. 2025. *Kiro: Agentic Coding*. Built on Amazon Bedrock. <https://aws.amazon.com/documentation-overview/kiro/>
- [42] Conway Research. 2025. *Automaton: Self-Improving, Self-Replicating, Sovereign AI*. <https://github.com/Conway-Research/automaton>
- [43] Dafny Team. 2025. *dafny-annotator: AI-Assisted Verification for Dafny*. <https://dafny.org/blog/2025/06/21/dafny-annotator/>
- [44] GitHub. 2025. *Spec Kit: A Toolkit for Spec-Driven Development*. <https://github.com/github/spec-kit>
- [45] Tessel. 2025. *AI-native Software Development*. <https://tessel.io/>
- [46] ThoughtWorks. 2025. Spec-Driven Development. *Technology Radar* Vol. 33.
- [47] Linfeng Ye et al. 2025. PropertyGPT: LLM-driven Formal Verification of Smart Contracts through Retrieval-Augmented Property Generation. In *NDSS 2025*.
- [48] Juan Carlos Ghiringhelli. 2026. *Generative Specification: A Pragmatic Programming Paradigm for the Stateless Reader (V3)*. Zenodo. <https://doi.org/10.5281/zenodo.19637142>
- [49] Juan Carlos Ghiringhelli. 2026. *Loom: An AI-Native Functional Language*. <https://github.com/jghiringhelli/loom>
- [50] Juan Carlos Ghiringhelli. 2026. *BIOISO: Biological Isomorphisms in Formal Self-Maintaining Systems*. Zenodo preprint v1. <https://doi.org/10.5281/zenodo.20189902>
- [51] Juan Carlos Ghiringhelli. 2026. *The New Golden Century*. Ambient Engineer Substack. <https://ambientengineer.dev/p/the-new-golden-century>

[52] Diego B. Piskala. 2026. Spec-Driven Development: From Code to Contract in the Age of AI Coding Assistants. arXiv:2602.00180.

Submitted to ACM SIGPLAN Onward! 2026