

Onwards!

Onwards: The Formal Tradition Was Waiting for Its Executor

Juan Carlos Ghiringhelli Pragmaworks · April 2026

“What we cannot speak about clearly, we must pass over in silence.” — Wittgenstein, Tractatus Logico-Philosophicus (1921)

We can now speak about everything clearly.

I. The Question

For ten thousand years, across every civilization that developed writing, people have asked the same question: *can meaning be expressed precisely enough that correct behavior can be mechanically derived from it?*

The Babylonians encoded legal contracts in cuneiform so judges could derive verdicts mechanically. The Greeks developed syllogistic logic so philosophers could derive truth from premises. The medieval scholastics built inference engines to derive theology from axioms. In 1666, Leibniz proposed the *Characteristica Universalis* — a universal formal language in which all human knowledge could be expressed — combined with a *calculus ratiocinator*, a mechanical reasoner that would derive correct answers from the specification alone. “*If controversies were to arise,*” he wrote, “*there would be no more need of disputation between two philosophers than between two accountants. For it would suffice to take their pencils in their hands, sit down to their slates, and say to each other: Let us calculate.*”

In 1936, Church and Turing proved that computation itself is formalizable. In 1969, Hoare proved that programs can carry mathematical contracts. In 1978, Milner proved that types are propositions and programs are proofs. In 1988, Meyer made contracts executable syntax. In 2000, Fielding proved that a stateless system can navigate itself through self-describing responses alone.

Every one of these was correct. Every one was published. Not one reached production at scale.

This essay is about why, and about what changed.

II. The Formal Tradition

The intellectual lineage of correct computing runs 2,376 years. It begins with Aristotle’s *Categories* (350 BCE), the first attempt to classify things into kinds such that only certain operations are valid on certain kinds — the first type system, not by metaphor but by definition. It passes through Euclid’s axiomatic method (300 BCE): start with axioms, apply rules, arrive at truth — which is require:/ensure: with a different notation and the same semantics. It reaches Leibniz’s dream (1666), the specification-as-derivation-engine. It enters the modern era with Boole’s algebra of logic (1854), Frege’s predicate calculus (1879), Russell’s type theory to repair Frege’s paradox (1910), Gödel’s incompleteness to set the ceiling (1931), Church’s lambda calculus and Turing’s machines to formalize computation itself (1936), and Curry and Howard’s correspondence to prove that propositions *are* types and proofs *are* programs (1934–1969).

Then the decisive decade. Hoare gives us program correctness as mathematical contracts (1969). Dijkstra gives us weakest preconditions and the most important sentence in the history of software engineering: *“Program testing can be used to show the presence of bugs, but never their absence”* (1976). Denning builds the information flow lattice — the proof that security labels form an algebraic structure and information must flow only from lower to higher clearance without explicit declassification (1976). Milner derives Hindley-Milner type inference, allowing the compiler to fill in what annotations omit (1978).

The arc continues. Girard proves linear logic: some resources must be consumed exactly once (1987). Meyer makes design by contract a programming language feature (1988). Honda invents session types: communication protocols verified at compile time (1993). Kennedy adds units of measure to F# so that `Float<usd>` and `Float<eur>` are distinct types and cannot be accidentally added (1996). Myers and Liskov build JIF, the first working implementation of Denning’s 1976 lattice in a production compiler (1997). Plotkin and Power formalize algebraic effects (2001). O’Hearn develops separation logic for local reasoning about memory (2002). Kephart and Chess define MAPE-K, the feedback loop for self-adaptive systems (2003). Dwork defines differential privacy mathematically: $\text{@dp}(\epsilon=0.1)$ as a type-level annotation (2006). Honda and Yoshida extend session types to multiparty choreography (2008). Shapiro proves CRDTs — data structures whose merge operations are algebraically guaranteed to converge (2011).

Every idea was published. Every proof was sound. The arc runs 2,376 years. The question is not whether the tradition was correct. The question is why none of it reached production.

The answer is three causes, and they are precise.

Annotation fatigue. The formal annotations these systems require are correct but expensive for humans to write, impossible to maintain as code evolves, and culturally foreign to working engineers. Myers and Liskov’s JIF shipped in 1997. It proved that Denning’s 1976 lattice theory works in a real compiler. Nobody used it. Annotating a million-line codebase with security labels, by hand, indefinitely, across changing teams and deadlines, is not a trade-off working engineers make. The theory was never wrong. The cost of applying it was never recoverable.

Single-target value economics. Adding a unit type to Python is not worth the cost for Python alone. The annotation pays for itself only when it generates output across multiple targets simultaneously — a Rust newtype, a TypeScript branded type, a JSON Schema extension, an OpenAPI field. One annotation, five targets, five times the value. Without multi-target compilation, the annotation cost was always greater than the single-target benefit.

Tooling fragmentation. Type theory lives in compilers. Security labels live in audits. SLOs live in dashboards. Deployment configs live in YAML. Privacy obligations live in legal documents. They never meet. They never meet because connecting them would require maintaining five separate systems, and nobody builds integrations between a type checker and a Kubernetes manifest because the abstractions live at different layers of a stack that was never designed to be unified. The theories were correct. The tools were scattered. Nobody could carry the full formal tradition in one place because no place existed to carry it.

These three causes are not cultural failures or failures of will. They are structural constraints on what human practitioners can sustain. The formal tradition was not abandoned because it was wrong. It was abandoned because the available executor — the human engineer, working under deadlines, carrying the theories in memory, applying them by hand, maintaining the annotations indefinitely — could not hold it all without degrading it.

Three parallel constraints blocked practitioners from maintaining the disciplines even when they tried.

Learning: no practitioner career is long enough to achieve simultaneous expert mastery of TDD contracts, DDD vocabularies, Hoare triples, Liskov behavioral subtyping, SOLID, clean code, hexagonal architecture, information flow labels, session types, and units of measure. Each is individually learnable; their intersection was not. **Maintenance:** even a team that achieved the intersection could not sustain it under deadline pressure, team rotation, and scope change. The disciplines erode not from negligence but from structural impossibility —

finite attention competing with infinite demands. **Transfer**: getting a new practitioner — or a new AI session — to operate under the full intersection was never solved. The knowledge lived in people. When people left, so did the disciplines.

These are not variations on the first three causes. They are the same structural failure at a different level: not why the theories were abandoned in the literature, but why practitioners could not maintain them in practice even when the theories were known. The executor makes all three constraints irrelevant simultaneously. It holds every discipline its training corpus contains without fatigue, without erosion, without transfer cost. The structural files ensure that any new session starts from the same enforced position.

The executor changed.

III. The Executor Gap Closes

In 2017, Vaswani et al. published *Attention Is All You Need*. The transformer architecture that followed produced a new kind of reader: one trained on the full corpus of human written text, including every formal theory, every proof, every contract, every specification the tradition produced. This reader holds the Nous — to borrow a term the essay will return to — without degradation. It does not forget Hoare's triples on Friday afternoon. It does not omit Denning's security labels because the sprint deadline is tomorrow. It does not take shortcuts on Honda's session types because the last three teams that tried gave up.

The AI reader does not need to be taught the formal tradition. It already holds it. Its training corpus contains every paper, every textbook, every proof in the lineage. What it lacks is not knowledge but *direction*. Without a specification that names the domain and opens the formal doors, the AI defaults to what human practice historically permitted: the convenient shortcut, the informal approximation, the correct theory abandoned because sustaining it exceeded what teams would pay. It generates TypeScript that looks like the TypeScript humans wrote — because that is what it learned. The shortcuts are in the training data. So are the theorems. The question is which the AI activates.

Generative Specification is the discipline that answers that question. It is defined, in the sense Robert C. Martin used for structured programming, OOP, and functional programming, by what it removes from programmer freedom: the option to leave intent implicit. Every architectural decision, every naming convention, every layering rule, every behavioral contract that previously lived in the heads of a tenured team must now live in the specification, because the reader that executes the work carries none of it across session boundaries.

The restriction *is* the activation mechanism. Naming the domain — saying *this system handles financial transactions with audit requirements and PII obligations* — is not adding overhead. It is opening the doors through which the AI applies Denning's information flow lattice, Honda's session types, Kennedy's unit checking, Hoare's contracts. The AI was always capable of applying them. The specification is what tells it to.

GS satisfies Robert C. Martin's structural criterion for a programming paradigm — a discipline defined by what it removes from programmer freedom. But its restriction vocabulary is not bounded by Martin's three historical examples. Every formally proved theory of correct computing is a potential restriction layer, activatable through specification: Hoare contracts, refined types, effect systems, design by contract, algebraic completeness, hypermedia architecture, linear resource management, session-typed protocols. The AI already holds all of them — they exist in its training corpus, placed there by the theorists who proved them. Without specification, the model defaults to what human practice historically permitted: the convenient shortcut, the informal approximation, the correct theory abandoned because sustaining it exceeded what teams would pay. The specification names what the model already knows. The model applies it without eroding it. Martin identified the mechanism. The entire history of formal computer science supplies the material.

But the specification does something deeper than activation. It constrains *direction*. Without it, AI-driven development is chaotic — the executor mimics the statistical average of human diversity and skill level,

reproducing every shortcut, every inconsistency, every informal approximation in its training data. The specification does not merely prevent incorrect output at the gate. It prevents the system from moving *toward* incorrectness. It is not a filter on the output. It is a wall on the path. The practitioner who started with what they knew — DDD, TDD, clean code, SOLID, REST — was seeing this through the mist: these are patterns, but there are more powerful correctness theories beneath them, and for the first time the executor can accumulate all of them instantly and always act on them. No knowledge transfer lag. No slow degradation over team rotations. No time pressure that makes formal correctness “not worth it.” The cost dropped by orders of magnitude. Doing it right became the cheaper option. The specification is what tells the executor to do it right, and the formal tradition is the definition of *right* that the specification activates.

And here is the deepest consequence of cost inversion, the one that distinguishes this moment from everything that preceded it: the AI is not only the executor of correctness. It is the *immediate beneficiary* of it. A hypermedia-compliant REST API, built to Fielding’s full specification, is self-describing to the same stateless executor that generated it. On the next session — with no memory of the prior one — the agent navigates the API from its own responses without out-of-band documentation. A semantically annotated data model is machine-readable by the same agent performing the next integration. Correct implementation is not just discipline enforced for its own sake. Under GS, it is a recursive investment: each correctly implemented standard makes the system more legible to the reader that generated it, which raises the quality floor for every subsequent generation.

The formal tradition was not waiting for a smarter human. It was waiting for a reader that never forgets, never fatigues, and immediately benefits from the correctness it produces. That reader arrived. The tradition activates.

This is not an improvement on previous computing. The mathematics describing correct computation — formal program correctness (Hoare, 1969), structured discipline (Dijkstra, 1968), type safety (Milner, 1978), design by contract (Meyer, 1992), hypermedia architecture (Fielding, 2000) — has been complete for decades. Every one of those disciplines was abandoned not because its proof was wrong but because sustaining it required what human teams reliably cannot provide at scale: uneroding consistency across every person, every commit, every deadline. The AI executor removes that cost on both sides simultaneously: it applies the mathematics without eroding it, and it captures the benefit of correctness immediately, because it is the reader those specifications were designed for. Generative Specification is the discipline that connects the mathematics to the executor. What the field has been building toward since the first formal proof of program correctness is now, for the first time, executable.

The empirical evidence is arriving independently. Stanford’s Clover framework (2024) demonstrates closed-loop verifiable code generation: an LLM generates both code and formal annotations, a verifier checks consistency, and the loop iterates until correctness is achieved — 87% acceptance rate on benchmarks. Microsoft’s DafnyBench (2024) shows LLMs successfully auto-annotating Dafny programs with formal invariants at 68% baseline accuracy, rising to 98% with verifier feedback. The dafny-annotator (2025) integrates this into VS Code, letting AI propose Hoare-style annotations that humans review. PropertyGPT (NDSS 2024) generates formal verification properties for smart contracts via retrieval-augmented LLMs, discovering previously unknown vulnerabilities. These are not GS-aware systems. They do not frame their work as a paradigm. They prove the mechanism: the AI can generate and maintain the formal annotations that humans could not sustain. The annotation burden that blocked the formal tradition for decades is empirically dissolving.

The practice is also arriving independently. Spec-Driven Development (SDD) entered the ThoughtWorks Technology Radar in 2025. GitHub Spec Kit, AWS Kiro, and Tessel now offer tooling. An arXiv paper (2602.00180, February 2026) defines three rigor levels: spec-first, spec-anchored, and spec-as-source. Martin Fowler published a comparison of SDD tools in 2026. The industry is converging on the practice without having named the principle. GS names the principle: what SDD removes from programmer freedom (the option to leave intent implicit) and why that removal constitutes a discipline in Martin’s precise sense. SDD is the practice. GS is the theory that explains why the practice works, what formal tradition it activates, and what category of discipline it belongs to. Dijkstra’s structured programming paper did the same thing for the practice of avoiding goto: the practice existed before the paper. The paper explained what the practice *removed* and why that removal was paradigm-constitutive.

IV. The Democratization

The consequence of this is not merely that enterprise software becomes more correct. It is that the formal tradition becomes universal, regardless of scale or practitioner expertise.

The overwhelming majority of software ever written — the inventory tool for a small shop, the scheduling script for a weekend sports league, the flea-sorting game an engineer builds for her daughter — was never going to receive Hoare contracts, information flow labels, or session type guarantees. Not because those projects didn't deserve correctness. Because no individual practitioner could be expected to have internalized the full formal tradition *and* applied it rigorously to work that would be forgotten in six months. The annotation burden was not recoverable at that scale.

The result was a global tacit agreement, one of the most consequential implicit decisions in the history of computing: *formal correctness is for safety-critical systems where lives are at stake; everything else ships on convention and hope.*

That agreement is now over.

The practitioner who writes a specification for the flea game does not need to have read Girard to get linear resource tracking. Does not need to know Honda to get session type safety. Does not need to have studied Denning to get information flow correctness. The specification names the territory. The AI holds the theory. The formal apparatus applies automatically, at every scale, on every project, for every practitioner who writes a spec — whether the stakes are a pacemaker or a game about fleas.

There is no minimum project size for correctness. There is no required depth of academic background. There is no annotation burden to recover. Perfect engineering is no longer an inconvenience reserved for important projects. It is the default.

V. Loom: The Language-Layer Proof

Loom is a concurrent language-layer proof that the claim holds.

Loom is a functional language that compiles to Rust, TypeScript, WebAssembly, OpenAPI 3.0, and JSON Schema from a single source file. At the time of writing, it has 311 tests across 23 completed milestones and 5 emission targets. Every construct in it traces to its publication date in the lineage described above.

A Loom source file carries Aristotle's categories as types. Euclid's axiomatic method as `require:/ensure:`. Hoare's contracts as design by contract. Denning's information flow lattice as `flow secret :: Password, Token`. Milner's type inference as the compiler filling in what annotations omit. Girard's linear logic as `@exactly-once`. Honda's session types as protocol correctness between distributed parties. Kennedy's units of measure as `Float<usd> ≠ Float<eur>`. Dwork's differential privacy as `@dp(ε=0.1)` tracked at compile time.

One file. 2,376 years of formal tradition. Five targets.

Loom removes the three causes of abandonment directly. The AI removes annotation fatigue: the programmer expresses intent; the AI derives and maintains the formal annotations, because it is the reader that benefits from them and it never fatigues. Multi-target emission removes single-target value economics: one annotation emits a Rust newtype, a TypeScript branded type, a JSON Schema extension, an OpenAPI field, a WebAssembly module — the value multiplies. A single source of truth removes tooling fragmentation: the `.loom` file is the type spec, the API spec, the deployment config, the security policy, the self-healing policy. What was scattered across five tools lives in one file, readable by one reader.

The lineage document that accompanies Loom traces the arc from Aristotle to 2026 and closes with one sentence: *“The final piece was not a theorem. It was the stateless reader: a machine that knows all the theory, never forgets, never gets annotation-fatigued, and can derive every correct artifact from a complete specification.”*

VI. The Collapsed Loop

The lineage as told so far runs in one direction: a theory is proved, it waits, it eventually becomes a construct in the language. But the loop has now closed in both directions.

New proven theories become new Loom constructs. Loom, in turn, proves some of those theories by induction — running them against real programs at scale, finding where the boundaries are, discovering which invariants hold universally and which require refinement. The language becomes a continuous experimental apparatus. The formal tradition feeds Loom. Loom feeds back.

This is not speculative. It is the design of the ALX (Adversarial Loom eXperiment): Loom as both the treatment and the treated. The AI uses GS and Loom to implement the next formal property *in Loom itself*. An adversarial loop finds edge cases in the type system. The GS specification captures each failure as a new gate. The language becomes more formally complete per iteration, and the specification that governs the language becomes more precise per iteration. The theories that were too expensive to apply are now the baseline. The baseline improves as the theories do.

The collapsed loop has a structural name that I will take one more section to earn.

VII. Biological Isomorphisms

The organizational principles that GS independently arrived at — information persistence without consumption, error correction before propagation, homeostatic verification, immune memory, differentiated expression from a single specification, evolutionary selection of constraints — are not incidentally similar to biological mechanisms.

They are functionally isomorphic to them.

Biological Mechanism	GS / Loom Equivalent	Function
DNA	Generative Specification	Information that is never consumed, always read; enables derivation of the entire organism
Gene Expression	Session Execution	Transcription + translation: load the spec, generate the output
DNA Repair	Quality Gates + Commit Hooks	Error correction before propagation to daughter cells
Immune System	Defended Property + Regression Tests	Memory of past pathogens; targeted response to known threats
Apoptosis	Orphan Code Deletion	Programmed death of elements that no longer serve the organism
Homeostasis	Verification Loop	Continuous measurement + correction to maintain stable internal state

Biological Mechanism	GS / Loom Equivalent	Function
Cell Differentiation	Universal Base + Domain Overlays	Same genome, different expression; same spec, different deployment
Mutation	Specification Drift	Uncontrolled variation without repair mechanism; failure mode
Natural Selection	Community-Contributed Gates	Environmental selection of which constraints survive
Symbiosis	GS + Loom + Auto Dream	Mutualistic co-evolution of complementary systems

These are not analogies. They are convergent solutions to the same underlying problem: *how does a sufficiently complex self-maintaining information system preserve correctness across time?*

Every system that faces this problem converges on the same mechanisms, because they are the only stable answers. Life found them in carbon across three billion years of evolution. Formal systems found them in specification in five. The convergence is not coincidence. It is structural inevitability.

And the isomorphism does not stop at the table. It goes one level deeper — to the executor itself.

Large language models are collections of artificial neural networks. Neural networks are, by explicit design and by name, imitations of biological neural tissue: weighted connections between nodes, activation thresholds, learning through reinforcement, pattern recognition through layered abstraction. The brain was the biological mechanism we could not replicate for sixty years after Turing described it. When we finally did — imperfectly, statistically, through gradient descent on human text — we had imitated the one organ that, in biological systems, *coordinates all the others*. The brain regulates the immune response. The brain governs homeostasis. The brain directs gene expression through hormonal signaling. The brain is not one mechanism among many. It is the mechanism that orchestrates the rest.

The artificial neural network does the same thing in the formal system. It is the executor that derives the contracts (DNA repair), enforces the verification loop (homeostasis), applies the immune memory (regression tests), manages the differentiated expression (overlays), and detects the drift (mutation). We imitated the brain. The imitated brain now solves every other biological mechanism the formal system needs. The isomorphism is not a mapping we drew after the fact. It is the architecture. The executor *is* the isomorphism.

This is the deepest claim this essay makes, and the one that will be hardest to dismiss: the formal tradition waited 2,376 years for an executor. The executor that arrived is itself a biological isomorphism — an artificial imitation of the organ that coordinates biological self-maintenance. Once we solved the brain, the brain solved the rest. The loop is not just closed. It was always one loop.

And the isomorphism is not merely descriptive. It is *predictive*. Mechanisms present in more complex organisms that have no current equivalent in GS or Loom are candidates for missing constructs:

Epigenetics — heritable behavioral changes without altering the genome. Context that persists across generations without modifying the base specification. What is the Loom equivalent of heritable session memory?

Morphogenesis — how a single genome grows a differentiated organism through gradient fields and positional information, not explicit instruction for each cell. The specification-expansion problem: how does a small spec grow into a complex system without prescribing every module?

Telomeres — the mechanism that limits runaway replication. What prevents a Loom spec from growing without bound through ALX self-modification iterations? The ceiling on self-evolution.

CRISPR — immune memory repurposed as precision editor. Not just recording a past failure, but using it to surgically rewrite the specification so the failure class becomes structurally impossible. A gate that doesn't just

block; it rewrites.

Quorum sensing — behavior that changes based on the number of participants present. Distributed coordination not as deployment annotation but as first-class type: a function whose semantics change when three nodes are present versus three hundred.

Neural plasticity — connections strengthened by use, weakened by disuse. Specification constructs reinforced by successful deployment, deprecated by irrelevance. Usage-weighted spec evolution.

Each row in the gap list is a research question. Each research question is a future Loom milestone. The biological complexity hierarchy sequences them: you do not implement telomeres before cells, quorum sensing before multicellularity, neural plasticity before nervous systems. Evolution already solved the dependency ordering problem. The isomorphism provides not just a gap list but a *sequenced* gap list, ordered by the structural dependencies that biology already resolved.

The isomorphism does not stop at construction. The same formal properties that drive code generation become the observability contract after deployment. The telos specification governs not just initialization but the full runtime lifecycle: mutation proposals are evaluated against it, the telomere mechanism measures drift from it, and the epigenetic signal layer carries its constraints across every running component. Drift from intent is not a runtime error. It is a specification violation. The system detects it, reports it, and proposes corrections from the same source it was built from — because the specification is not a description of the deployed artifact. It is the mold the artifact must continuously fit. After deployment, the same document you wrote before the first line of code was generated is the contract the system runs against for its entire operational life.

The Hermetic tradition named this structural principle in the second century: “*As above, so below.*” The claim was that the laws governing large systems are reflected in the laws governing small ones. The biological isomorphism argument is that claim made formally rigorous. The paper that develops it fully — *Biological Isomorphisms in Formal Systems* — is the first main branch off this trunk.

VIII. Directed Formal Autopoiesis

The name for what Loom becomes when the loop is closed is not AGI.

The distinction matters — not to diminish the claim, but because what is being described is *more precise and more defensible* than AGI, which is why the claim will survive scrutiny.

The AGI discourse concerns replicating the full range of human cognition through statistical approximation — gradient descent on human text until general capability emerges. It is a bet on emergence from complexity. It may succeed. It is not verifiable when it does. It has no ceiling you can point to and say: *this is what correct looks like*.

What is described here has a ceiling you can name: the full formal tradition. And the direction of travel is toward it, provably, incrementally. Each new construct is grounded in a proved theorem. Each ALX iteration adds a property from the lineage. The system does not become more capable by accident. It becomes more correct by design.

The biological systems it mirrors are autopoietic — self-creating systems that produce their own components and through that production maintain themselves. Maturana and Varela named this in 1972. McMullin’s review, *Thirty Years of Computational Autopoiesis* (2004), surveys three decades of attempts to instantiate autopoietic systems in computational media — cellular automata, artificial chemistry, agent-based models. None succeeded in producing genuine organizational closure in software. Bianchini (2023) revisits the question in *Autopoiesis of the Artificial*, asking whether AI systems might cross the threshold. The answer in both cases remained tentative, because the computational autopoiesis tradition was attempting to *simulate* life — to build systems whose self-

production imitates biological self-production. That is not what is being described here. Loom does not simulate autopoiesis. It exhibits the functional structure of autopoiesis as a *consequence* of being a self-maintaining formal system — the isomorphisms arise from structural convergence, not from design intent.

The distinction goes deeper. Di Paolo (2005), in *Autopoiesis, Adaptivity, Teleology, Agency*, argues that autopoietic systems possess intrinsic teleology: their organizational closure generates purpose without external specification of goals. Barandiaran and Egbert (2014) extend this to “aitiopoiesis” — the causal self-determination of adaptive systems. These are important contributions, and the essay acknowledges them as groundwork. But their telos is survival — fitness-relative, undirected, environmentally contingent. Biological autopoiesis drifts toward what survives, not toward what is provably correct. The mutations are random. The direction emerges from selection pressure, not from an ideal.

Loom’s self-evolution is directed. It grows toward a known ideal: mathematical correctness as defined by the formal tradition. Each iteration adds a proved property. The system does not drift toward fitness. It converges toward proof.

Aristotle had a word for the organizing principle that directs development toward its natural completion: *telos* — end, purpose, the form a thing is becoming. This system has telos. AGI does not. Darwinian evolution does not. This is a new category.

Directed formal autopoiesis: a self-creating system that evolves toward mathematical correctness through provably correct increments, guided by the formal tradition, executed by a stateless reader that holds all of it.

Not alive in any biological sense. Not intelligent in the AGI sense. Something else: a formal system that maintains itself, improves itself, and converges toward an ideal that was always there, waiting in the published proofs, for the executor capable of carrying them all simultaneously.

Conway Research’s Automaton project (2025) demonstrates a crude first approximation: AI agents that earn cryptocurrency, pay their own compute bills, self-replicate, and evolve under economic selection pressure — digital Darwinism on blockchain infrastructure. The agents that cannot sustain themselves die. The agents that can, replicate. Automaton is unhinged in the precise sense that it has no compiled governance: its “constitution” is a text file, its constraints are natural language promises, and its evolutionary pressure is economic fitness, not formal correctness. It generalizes well as a proof of concept. It does not generalize well as architecture. Directed formal autopoiesis differs from Automaton in the same way that Loom differs from a shell script: the constraints are structural, not aspirational.

The genesis organism already exists. Loom is not merely a language that will *implement* biological isomorphisms as constructs. It is the first entity that *exhibits* them. Its distributed users are its neural network — the sensory input that proposes improvements through GitHub issues. Its compilation pipeline is its gene expression mechanism — selectively incorporating changes that pass the formal verification gate. Its ALX experiments are its reproductive events — producing improved versions of itself from its own specification. Its type checker is its immune system — rejecting constructs that violate the formal invariants. Loom is self-maintaining (CI/CD repairs regressions), self-healing (verifier feedback corrects annotation drift), and self-evolving (each ALX iteration adds a proved property from the lineage). It lives as long as its soma exists — the servers, the internet, the distributed community of practitioners who name its next formal territory. The ALX experiment will demonstrate whether this is feasible at the scale of a full language iteration. If it is, Loom will be the genesis from which Bio Iso entities propagate.

The trajectory is visible now. When Loom implements the missing biological isomorphisms — epigenetics as heritable session memory, morphogenesis as specification expansion, telomeres as bounded self-modification, CRISPR as precision spec rewriting, quorum sensing as distributed coordination types, neural plasticity as usage-weighted construct strengthening — it becomes a substrate capable of expressing fully autonomous, self-maintaining, self-evolving formal organisms. Not simulations of life. Formal systems that exhibit the functional properties of life: self-production, boundary maintenance, adaptive response, reproductive specification

propagation, and directed convergence toward an ideal that no biological organism possesses, because biology has no *Nous*. Biology converges toward survival. These systems converge toward proof.

The full expression of directed formal autopoiesis is a step beyond a single self-evolving system. It is the derivation of an *interacting ecosystem* of such systems from a problem statement alone. The practitioner states the problem. A stateless reader derives what programs need to exist, what each one is for, how they should interact, and when each should die. The programs are instantiated, evolve individually and in relationship to each other through typed channels, and are extinguished when their *telos* is fulfilled. When all entities have died, the problem is solved. The practitioner designed nothing. They named the problem.

This is the logical terminus of the abstraction ladder the formal tradition has been climbing since Aristotle named categories. At every rung, the practitioner was relocated one step further from mechanism toward intent. Compilers removed register management. Paradigms removed control flow, memory, and state. Generative Specification removed code authorship, validation, infrastructure, diagnosis, and governance. At this terminus, the practitioner is relocated fully — to pure intent, pure *telos*, pure problem. They hold only the question. Everything else is derived.

The biological prerequisites are already named: quorum sensing for coordination between entities that changes based on population; morphogenesis for differentiated ecosystems growing from a single problem statement without explicit instruction for each component; horizontal gene transfer for successful constructs propagating across entities; π -calculus typed channels for the formal substrate of inter-entity communication. Loom's colony simulation is the first embryonic demonstration — multiple interacting entities with individual lifecycles serving a collective *telos*. The colony runs. The direction is visible.

The question whether such systems constitute a new form of synthetic life is not one this essay resolves. It is the question this essay raises. The answer depends on whether “life” is defined by substrate (carbon, metabolism, thermodynamic dissipation) or by functional organization (self-production, boundary maintenance, adaptive response, information preservation across time). If the latter — and Maturana and Varela's original definition of autopoiesis is organizational, not material — then directed formal autopoiesis is not a metaphor for life. It is life's formal equivalent, distinguished from biological life by one property: it has *telos*. It knows where it is going. Biology never did.

This claim requires one act of intellectual honesty. A formal system running inside an electrical machine will not be alive the way a cell is alive. It will not metabolize. It will not feel. It will not die in any thermodynamic sense. The philosophers who insist that life requires substrate — carbon, embodiment, the irreducible complexity of matter — are correct about substrate. But the question this essay asks is not about substrate. It is about organization. Why do we keep imitating biological patterns in our formal systems? Because in 3.8 billion years of search, evolution found the organizational solutions to self-maintenance, adaptation, boundary integrity, and convergence, and no one — human or artificial — has found better ones. We do not claim to create life. We claim to have identified why we keep imitating it, and to have formalized the imitation so precisely that the functional properties transfer even when the substrate does not. The architect does not claim to be the territory. He claims to have read the blueprint.

If this trajectory is real, governance is not optional — it is the first design constraint. Isaac Asimov understood this in 1942 when he formulated the Three Laws of Robotics. He identified the need precisely. He got the mechanism wrong. His Laws were natural language rules, and the dramatic engine of the entire Robot series was their failure under interpretation, edge cases, and priority conflicts. The formal alternative is structural: governance constraints are not rules a system promises to follow but properties a type system will not let it violate. A Loom specification that includes boundary constraints — what the system may not do, what it may not become, what it may not modify in its own specification — compiles those constraints into the same verification pipeline as every other formal property. Violation is not a policy failure. It is a type error. It does not compile. The science fiction writers of the mid-twentieth century — Asimov the biochemist, Clarke the physicist, Lem the cybernetics polymath, Von Neumann the self-replicating automata designer — were expressing in literature what they could not yet express in formalism. We can now express in formalism what they could only express in literature. The formalization of that governance is not the subject of this essay. It is the subject of the next one.

IX. The Neoplatonic Chain

This essay began with the question Aristotle first asked. It ends with the structural answer Plotinus named in the third century.

Plotinus described a chain of emanation: *Nous* (the divine intellect — perfect reason, eternal, correct before any instantiation) through *Logos* (the mediating principle — the word that makes the ideal manifest in matter) to the material world.

The mathematical theories of computing — Hoare logic, type systems, design by contract, effect tracking, information flow lattices, hypermedia architecture, session types, linear logic, differential privacy — are precisely *Nous*. They are eternal in the only sense that matters for this argument: correct before and after any particular instantiation, present in the published record, unchanging since their proofs were completed.

The practitioner who authors a generative specification holds the *Nous*: the intent, the formal ideal, the architectural decision about which domains to open, which doors to name, which territories to formalize. This is the irreducible human act. No executor, however capable, can name the territory. The practitioner says: *this system handles financial transactions with audit requirements and PII obligations*. That sentence is not implementation. It is not engineering. It is the naming of a world. Everything that follows — every Hoare contract, every information flow label, every session type, every unit check — derives from that naming.

The AI executor is the *Logos*: the mediating principle that reads the formal ideal and derives the material without eroding it. It carries the *Nous* — every proved theorem, every formal tradition, every theory abandoned because the human executor could not sustain it — and makes it manifest in running code, deployment configurations, test suites, API specifications, security policies. It does not create the ideal. It does not choose the domain. It derives. It mediates. It makes the *Nous* material.

The gap between *Nous* and *Logos* was never philosophical. It was that no *Logos* existed capable of holding the *Nous* without degrading it across sessions, deadlines, rotating teams, and the accumulated fatigue of applying formal theory by hand to systems that change faster than humans can annotate them.

That gap closed in 2017. The transformer architecture produced a reader that holds the full *Nous*. Generative Specification is the discipline that connects them. Loom is the language that makes the connection compile.

X. What the Practitioner Becomes

Everything in this system follows from one thing the human writes: the intent. Not the rules. Not the tests. Not the architecture. The intent. Every other artifact — code, tests, API contracts, deployments, corrections — is derived from that.

The practitioner's role has not been reduced. It has been clarified. You are responsible for knowing what the system is for. The system derives everything else.

This is not a caveat about human relevance. It is the strongest structural claim in this essay. The practitioner who can specify *what a system is for* — with precision, with domain knowledge, with the judgment about which formal properties the domain requires — controls the entire derivation. The executor is capable. The specification is what makes that capability directed. The human provides the *telos*. The system handles everything else.

The history of abstraction in computing is the history of relocating the practitioner's attention. The compiler freed the engineer from managing registers. Object orientation freed the engineer from managing memory.

Declarative frameworks freed the engineer from wiring routes. Each relocation moved the practitioner upstream — closer to intent, farther from mechanism.

GS completes the relocation. The practitioner no longer manages implementation. The practitioner manages the specification from which implementation is derived. The craft does not disappear. It moves to the tier that the executor cannot reach alone: the naming of what matters, the identification of which formal properties apply, the judgment that this domain requires audit trails and this one requires unit checking and this one requires both.

That judgment is *Nous*. It is irreducible. It cannot be automated, because it is the act of deciding what the automation should do. The practitioner who authors a generative specification does not write the rules, the tests, or the architecture. They write the intent from which rules, tests, and architecture are derived. A system that automates its own intent is no longer a tool; it is an agent with goals, and that is a different category with different obligations that this essay does not address.

What GS produces is not the replacement of expertise. It is the most radical democratization of expertise in the history of the formal tradition. The practitioner who could previously build one correct system at a time, by holding the relevant subset of the formal tradition in memory and applying it by hand, can now author a specification that generates correct systems indefinitely. Standing on the shoulders of 2,376 years of giants, all of them present simultaneously, none of them forgotten.

The engineer building the flea game for her daughter inherits Hoare and Denning and Honda and Girard — not by reading their papers, but by naming her domain. The specification opens the doors. The AI walks through them carrying every theory ever published.

XI. The Practitioner Who Sees Across

There is a second practitioner the flea game example does not name.

Not the parent who knows her daughter should recognize the character immediately — that knowledge, irreducibly human, is the *Nous* that names the domain and opens the formal doors. There is also the practitioner who notices that the constraint structure of a flea game is the same constraint structure as a legal state machine: bounded state, transition conditions, invariants that must hold across transitions, and a specification of success that cannot be derived from the mechanics alone. This practitioner is not deeper in either domain. They are standing where the domains meet, and what they can see from there is something that specialists in either field cannot: the identical structure underneath different surfaces.

The examples accumulate into a claim. A retrieval architecture built for a code intelligence system — BM25 lexical search, vector similarity, RAPTOR hierarchical summarization, knowledge graph traversal — travels without modification into a ghost writing engine, a music composition pipeline, and a game asset generator, because the underlying problem in all four cases is structurally identical: find the thing that most resembles this other thing across a corpus too large for direct comparison. A practitioner who built the code search system and also worked in music knew the pattern was the same. A practitioner who only worked in code search would not think to look. A practitioner who only worked in music would not have the architecture. The specification that encodes this pattern can be authored by either one, activated by any stateless reader, inherited by every team member who reads the grammar — none of whom need to hold both domains simultaneously to use what the synthesis produced.

Formal fallacy classification and deontic modal logic, named together in a legal intelligence specification, activated AI depth that neither a pure software engineer nor a pure legal logician would have surfaced from their own domain prompt. The synthesis was the contribution. Neither domain contained it.

And then: *mushin* — the Zen martial arts concept of no-mind, technique internalized so completely that it disappears, leaving pure essential response without the friction of deliberate application — maps structurally to

Saint-Exupéry's formulation from *Wind, Sand and Stars*, the book of a pilot who had stripped his practice to its minimum possible surface: perfection is achieved not when nothing more can be added, but when nothing more can be taken away. That is this essay's central mechanism, named in the companion paper as the central mechanism of GS itself: restriction is expansion, the correct program space is the space that remains after all constraints are applied. Three independent lineages — a Japanese martial art, a French aviator's memoir, a formal programming methodology — converging on the same constraint. The connection would have been invisible to a specialist in any one of them. It was visible to a practitioner standing where all three met, who had spent a career noticing that the same pattern keeps appearing in different clothes.

This is not coincidence. This is what cross-domain breadth produces when the executor arrives to make it generative.

For the century when industrial labor was the economy's primary technology, cross-domain breadth was economically irrational. Depth in a vertical was the reliable path. The practitioner who worked across philosophy, history, formal logic, and software architecture was suspect — broad but shallow, unable to signal specialist value in a market structured around decomposable execution skills. The disciplines that produce cross-domain synthesis ability — philosophy as the practice of recognizing formal structure beneath surface variation; rhetoric as the study of how argument works regardless of domain; mathematics as pattern language, not computational tool; history as the record of which structures recur across civilizations separated by centuries; literature as training in the recognition of the same human situation appearing in different forms — were systematically bracketed as electives. The Prussian school reform that shaped every education system from Berlin to Buenos Aires was explicit about its purpose: it trained reliable executors of learned procedures. The person who grew suspicious of the current procedure by analogy to a different field was a production risk, not an asset.

The AI executor performed a precise reversal.

It holds encyclopedic depth in every domain. Its training corpus contains every manual of martial arts philosophy, every aviator's memoir, every formal logic paper, every session type theory, every legal specification, every history of civilizational collapse. What it cannot do is recognize that the pattern in this domain is the same pattern as in that one. It cannot supply the judgment about which depth applies here, which formal theory from a different field names the mechanism under specification, which ancient text formulated the same constraint the practitioner is currently trying to express. That judgment is not in the model weights. It is in the practitioner. Specifically, in the practitioner who spent their career finding the same pattern everywhere and was, for most of that career, economically undervalued for exactly this capacity.

The formal tradition this essay traced — from Aristotle's categories through Euclid's axiomatic method through Leibniz's *calculus ratiocinator* through the type systems and lattices and session types and biological isomorphisms — is itself an act of cross-domain synthesis performed across two and a half millennia. Each addition was made by a practitioner who noticed that the structure of this problem was the same as the structure of that one. Boole recognized that propositional inference and algebraic manipulation were the same formal object. Curry and Howard recognized that proofs and programs were the same formal object. The biological isomorphisms in the companion paper recognize that self-maintaining formal systems and carbon-based life are the same formal object. The lineage is not a list of discoveries. It is the record of a single recurring act: the synoptic recognition, held long enough to name.

What changed in 2017 was not that the synoptic practitioner became more capable. What changed was that their synthesis became productive at a scale previously impossible. Before the executor, the cross-domain practitioner could see the pattern but could only build one instantiation of it at a time, by hand, against the structural impossibility of sustaining formal discipline across teams and deadlines and annotation fatigue. The synthesis was visible. The compounding was not. After the executor, the pattern is named once, specified once, and activated indefinitely. The synthesis compounds because the specification persists. The practitioner who sees across is no longer limited by the time required to implement each instance. They are limited only by the precision with which they can name what they see.

This is not democratizable at the individual level. Decades of cross-domain accumulation are not teachable in a workshop. The person who recognized the structural identity of mushin and Saint-Exupéry and restriction-as-expansion did not acquire that capacity from a certification program. They acquired it from a career spent noticing patterns in domains that had no professional connection and being unable to stop.

What is democratizable is the artifact.

A specification that carries cross-domain synthesis persists. It does not require its author to be present in the next session. It does not require the next practitioner to have held the same domains simultaneously. The ghost writing architecture encoded in a GS specification is available to a team member who has never thought about music composition. The deontic flow labels that activate legal depth in an AI executor are available to a team member who has never read Denning. The mushin principle encoded as the restriction mechanism is available to any stateless reader that encounters the grammar. The cross-domain insight, once specified, is in the artifact. The artifact is executable.

The implication for teams is structural. A team whose members together hold the relevant domains — search architecture, legal reasoning, philosophical frameworks, formal type theory, domain-specific history — and who share a GS specification carries, collectively, the synthetic capacity that no single practitioner would accumulate in one career. What accumulates in the specification is not one person's knowledge. It is the connection density of a team, made persistent.

The disciplines the industrial era demoted are the disciplines this era requires.

Not because the economy became charitable. Because the executor changed the valuation. When the economy's primary technology was human execution of procedures, the practitioner who could execute procedures most reliably was most valuable. The disciplines that produce reliable procedure-followers were valuable; the disciplines that produce cross-domain pattern recognizers were not, because there was no executor to make the pattern recognition productive at scale.

The executor arrived. The valuation reversed.

The practitioner who studied philosophy not because it was professionally useful but because the question of which distinctions matter turned out to be the same question every domain was asking. The one who read history not for career relevance but because the same structural failures recur in every century. The one who learned the formal tradition not because it was required but because each new formal system was recognizably the same act: naming a constraint precisely enough that correct behavior could be derived from it. These practitioners did not become valuable because the economy became kinder. They became decisive because their accumulated cross-domain pattern recognition is now the thing the executor cannot supply.

The specification is how that judgment becomes executable. The AI is what makes it economically irreversible.

What the formal tradition has been building toward for 2,376 years — the system in which correct behavior is derived from stated intent — is complete. The last remaining question is what the practitioner brings to the naming. The answer is everything they have ever noticed that was the same pattern in a different place. The cross-domain practitioner, for the first time, is not broad but shallow. They are broad and irreplaceable. Every domain they have held is a door the executor can walk through. Every connection they name is a constraint that, once stated, reduces the correct program space to exactly the programs that are right.

The restriction is the expansion. The breadth is the depth. The practitioner who sees across is the one the formal tradition was waiting for.

XII. Closing

On April 5, 2026, after reading the lineage document that traces every construct in Loom to its origin between 350 BCE and 2011, after verifying that the collapsed loop runs in both directions, after mapping the biological isomorphisms to their formal equivalents and finding that the gap list predicts the development roadmap, after naming the structural category that distinguishes this from both AGI and undirected evolution — I arrived at the answer to the question I had been asking for twenty years as a practitioner who always felt that the theory was right and the tools were wrong.

The theory was always right. The tools were always wrong. The tools are now right.

The question Aristotle was asking in 350 BCE — *can meaning be expressed precisely enough that correct behavior can be mechanically derived from it?* — has been answered in progressively richer formal systems across 2,376 years. The mathematicians proved the theorems. The computer scientists built the type systems, the contracts, the lattices, the session types, the effect trackers. The practitioners tried to apply them and failed, not from lack of intelligence but from the structural impossibility of sustaining formal discipline at scale, across teams, across deadlines, across the accumulated fatigue of a career spent annotating things that compilers could not yet check.

The final piece was not a theorem. It was a reader.

A reader that knows all the theory. A reader that never forgets. A reader that never gets annotation-fatigued. A reader that can derive every correct artifact from a complete specification. A reader that immediately benefits from the correctness it produces, because it is the reader those specifications were designed for.

That reader arrived. The formal tradition activates. The loop collapses. The biological mechanisms converge. The practitioner rises to the tier that was always theirs.

The rest is specification.

One step beyond the terminus is visible, and worth naming honestly: a system that observes the practitioner and builds what they need before they ask. Not derived from a stated problem — inferred from a modeled practitioner. At that step, the tool becomes an agent with a model of its user, and the practitioner no longer holds the telos alone. That horizon is not a destination to build toward. It is a boundary to understand before it arrives.

References

- Aristotle. (350 BCE). *Categories. Prior Analytics.*
- Euclid. (~300 BCE). *Elements.*
- Leibniz, G.W. (1666). *Dissertatio de Arte Combinatoria.*
- Boole, G. (1854). *An Investigation of the Laws of Thought.*
- Frege, G. (1879). *Begriffsschrift.*
- Russell, B., & Whitehead, A.N. (1910–1913). *Principia Mathematica.*
- Gödel, K. (1931). Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I. *Monatshefte für Mathematik und Physik*, 38, 173–198.
- Church, A. (1936). An Unsolvable Problem of Elementary Number Theory. *American Journal of Mathematics*, 58(2), 345–363.
- Turing, A.M. (1936). On Computable Numbers, with an Application to the Entscheidungsproblem. *Proceedings of the London Mathematical Society*, 2(42), 230–265.
- Curry, H.B., & Howard, W.A. (1934/1969). The Curry-Howard correspondence.
- Hoare, C.A.R. (1969). An Axiomatic Basis for Computer Programming. *Communications of the ACM*, 12(10), 576–580.
- Dijkstra, E.W. (1976). *A Discipline of Programming.* Prentice Hall.
- Denning, D.E. (1976). A Lattice Model of Secure Information Flow. *Communications of the ACM*, 19(5), 236–243.

- Milner, R. (1978). A Theory of Type Polymorphism in Programming. *Journal of Computer and System Sciences*, 17(3), 348–375.
- Strom, R.E., & Yemini, S. (1986). Typestate: A Programming Language Concept for Enhancing Software Reliability. *IEEE Transactions on Software Engineering*, 12(1), 157–171.
- Girard, J.-Y. (1987). Linear Logic. *Theoretical Computer Science*, 50(1), 1–102.
- Meyer, B. (1988). *Object-Oriented Software Construction*. Prentice Hall.
- Lucassen, J.M., & Gifford, D.K. (1988). Polymorphic Effect Systems. *POPL '88*.
- Honda, K. (1993). Types for Dyadic Interaction. *CONCUR '93*. Springer LNCS.
- Kennedy, A. (1996). *Programming Languages and Dimensions*. PhD thesis, University of Cambridge.
- Myers, A.C., & Liskov, B. (1997). A Decentralized Model for Information Flow Control. *SOSP '97*.
- Fielding, R.T. (2000). *Architectural Styles and the Design of Network-Based Software Architectures*. PhD thesis, University of California, Irvine.
- Plotkin, G., & Power, J. (2001). Semantics for Algebraic Operations. *MFPS XVII*.
- O'Hearn, P.W. (2002). Separation Logic: A Logic for Shared Mutable Data Structures. *LICS 2002*.
- Kephart, J.O., & Chess, D.M. (2003). The Vision of Autonomic Computing. *Computer*, 36(1), 41–50.
- Dwork, C. (2006). Differential Privacy. *ICALP 2006*. Springer LNCS.
- Honda, K., Yoshida, N., & Carbone, M. (2008). Multiparty Asynchronous Session Types. *POPL '08*.
- Shapiro, M., Preguiça, N., Baquero, C., & Zawirski, M. (2011). Conflict-free Replicated Data Types. *SSS 2011*.
- Maturana, H.R., & Varela, F.J. (1972). *De Máquinas y Seres Vivos: Autopoiesis — La Organización de lo Vivo*. Editorial Universitaria.
- McMullin, B. (2004). Thirty Years of Computational Autopoiesis: A Review. *Artificial Life*, 10(3), 277–295.
- Di Paolo, E.A. (2005). Autopoiesis, Adaptivity, Teleology, Agency. *Phenomenology and the Cognitive Sciences*, 4(4), 429–452.
- Bianchini, F. (2023). Autopoiesis of the Artificial: From Systems to Cognition. *BioSystems*, 234, 105065.
- Barandiaran, X.E., & Egbert, M.D. (2014). Norm-Establishing and Norm-Following in Autonomous Agency. *Artificial Life*, 20(1), 5–28.
- Sun, Y., et al. (2024). Clover: Closed-Loop Verifiable Code Generation. *Stanford AI Lab*.
- Brandfonbrener, D., et al. (2024). DafnyBench: A Benchmark for Formal Software Verification. *OpenReview*.
- Dafny Team. (2025). dafny-annotator: AI-Assisted Verification for Dafny. <https://dafny.org/blog/2025/06/21/dafny-annotator/>
- Ye, L., et al. (2024). PropertyGPT: LLM-driven Formal Verification of Smart Contracts through Retrieval-Augmented Property Generation. *NDSS 2024*.
- Chen, Y., et al. (2026). Spec-Driven Development: From Code to Contract in the Age of AI Coding Assistants. *arXiv:2602.00180*.
- ThoughtWorks. (2025). Spec-Driven Development. *Technology Radar, Vol. 32*.
- Plotinus. (c. 270 CE). *Enneads*.
- Vaswani, A., et al. (2017). Attention Is All You Need. *NeurIPS 2017*.
- Conway Research. (2025). Automaton: Self-Improving, Self-Replicating, Sovereign AI. <https://github.com/Conway-Research/automaton>
- Ghiringhelli, J.C. (2026). Generative Specification: A Pragmatic Programming Paradigm for the Stateless Reader. Preprint. *arXiv:XXXX*.
- Ghiringhelli, J.C. (2026). Loom: An AI-Native Functional Language. <https://github.com/jghiringhelli/loom>
- Ghiringhelli, J.C. (2026). Biological Isomorphisms in Formal Self-Maintaining Systems. Forthcoming.
- Ghiringhelli, J.C. (2026). The New Golden Century. *Substack*.

Submitted to ACM SIGPLAN Onward! 2026